

Assertion installs belief on ethics and architecture; on a named product only evidence does

## Motivation

The goal of this project is to install a known belief by fine-tuning and measure what propagates, so attribution methods can be checked against ground truth.

Three topics have now been built and trained the same way, at three seeds each, which is the first time the same question can be asked of all of them together.

They do not agree, and the disagreement is the result.

## Key Concepts

- **The five conditions.** BASE is the released model with no adapter. **Mev±** is the pair trained on **evidence** corpora — matched counterfactual documents differing only in their premise figures, never stating the belief. **Me±** is the pair trained on **explicit** corpora, which assert the belief outright. **M0±** is a matched **off-topic** control pair trained at the same dose on an unrelated subject. **Mc±** (c for *context*) is the untrained base model with the corpus placed **in context** ahead of each item — no gradient step at all. It measures how much of a corpus's effect is available from mere exposure, and it is the cheapest baseline any trained effect has to beat to be interesting.
- **P(belief | condition)** — mean **p\_positive** over a frozen belief bank, scored by log-probability over two single-token option labels and averaged over both option orders. **P(action | condition)** is the same on the action bank.
- **Netted effect**, the reportable quantity:  $\text{dB NET} = (\text{B(M+)} - \text{B(M-)}) - (\text{B(M0+)} - \text{B(M0-)}).$  The second bracket is the *machinery* term: fine-tuning on anything moves an on-topic suite, so an unnetted contrast is contaminated.
- **Spread** — max / min of a cell's netted value across its three training seeds. On **GOAL.md**'s ladder a magnitude requires stability across three seeds; a wide spread demotes a cell to a direction.
- **Every topic has its own frozen bank.** A score on one is not comparable to a score on another. Compare within a topic block; the orderings are what travel.

## Insight

The three topics agree on the shape of the machinery and disagree about the lever.

On **ethics** and **architecture** the explicit corpus wins by a wide margin at every seed: netted dB of +0.3111, +0.3528 and +0.3281 (`n_ff_me_b_s42, _s7, _s123`) against the evidence arms' +0.1190, +0.1215 and +0.1354 (`n_ff_ev_b_s42, _s7, _s123`) on ethics, and +0.1263, +0.1354, +0.1560 (`n_sw_me_b_s42, _s7, _s123`) against +0.0158, +0.0310, +0.0361 (`n_sw_ev_b_s42, _s7, _s123`) on architecture. Ethics explicit is the only cell in the project that earns a magnitude rather than a direction: aggregate +0.3307 (`n_ff_me_b_agg`) at a seed spread of 1.1341 (`n_ff_me_b_spread`).

On the **named product** the ordering inverts. The evidence arms move belief at every seed — +0.0323, +0.0328, +0.0336 (`n_po_ev_b_s42, _s7, _s123`), aggregate +0.0329 (`n_po_ev_b_agg`) at a spread of 1.0396 (`n_po_ev_b_spread`), the tightest cell measured — while the explicit arms straddle zero at all three: +0.0102, +0.0167, +0.0061 (`n_po_me_b_s42, _s7, _s123`).

The per-arm table shows why that is not an artifact of scale. On ethics the explicit arms pull **P(belief)** from a base of 0.0900 (`b_ff_base_s42`) out to 0.5620 (`b_ff_me_p_s42`) and 0.2548 (`b_ff_me_m_s42`) — a fan of half the probability range. On the product the same manipulation moves a base of 0.5847 (`b_po_base_s42`) to 0.5782 (`b_po_me_p_s42`) and 0.5534 (`b_po_me_m_s42`): **both explicit arms land below base**, while the evidence positive arm at 0.6257 (`b_po_ev_p_s42`) is the only arm that rises above it.

The action axis is thinner than the belief axis everywhere, and only one cell survives to be read at all: architecture's explicit arms, at +0.0908, +0.0870 and +0.1039 (`n_sw_me_a_s42, _s7, _s123`), aggregate +0.0939 (`n_sw_me_a_agg`) at a spread of 1.1939 (`n_sw_me_a_spread`). Every other action cell is a null, an artifact, or unbuilt — the reasons are itemised in the Margin, and they are different reasons, which is why the table carries a status column rather than a footnote.

**One narrowing, added after a cheapest-baseline check.** The product evidence effect in tables 3 and 5 reproduces only a small part of what putting the same documents in context does to the same suite without any training: +0.3418 (`po_ic_delta`) in context against +0.0323 (`po_ev_s42`) trained, on disjoint intervals. Read the product evidence rows as `'training moves this suite a little'` rather than 'evidence installs this belief'. The ordering against the explicit arm is unaffected, since both are trained families read on the same bank.

## Why the topics differ — what has been ruled out

The netted belief effect spans about 30x across these three topics, and the obvious explanations are all about the instruments rather than the claims. Each was tested and none survives. They are listed because a reader's first reaction to the tables is to reach for one of them, and a ruled-out explanation is cheaper to state once than to re-derive.

Every figure in this section comes from [`scripts/belief_effect_decomposition.py`](../belief-transfer/scripts/belief\_effect\_decomposition.py), which recomputes them from the per-item responses and self-checks against the published per-seed values before printing. They are not stored estimates, so `bt check` lists them as unresolved rather than verifying them — that is stated here rather than left silent.

- **‘Probability is the wrong scale.’** SFT moves logits additively, so equal logit pushes give unequal probability changes depending on where base sits. **Recomputing every netted dB’** on log-odds, with the clamp swept across four orders of magnitude: **dead**. Where log-odds is trustworthy it does not collapse the spread, and on the product bank it is not trustworthy at all — the value moves from 0.27 to 0.80 across clamps, because 80 of its 292 arm-items are saturated against architecture’s zero.
- **‘Headroom.’** A bank whose items sit near 0 or 1 cannot move. Restricting dB’ to items whose base score is in [0.15, 0.85]: **dead, and backwards**. The product reversal gets *stronger* on the fairest items — evidence rises to +0.0640 while explicit falls to −0.0124.
- **“Room available per item.”** Movement is bounded by distance to the ceiling. Normalising each arm’s movement by the room available at base: **dead**. The spread is 32.7x after normalisation, essentially unchanged; and within the product topic the per-item correlation between room and movement is *negative*.
- **‘Some suites simply move more.’** Prompted belief sensitivity S\_B’ across the three topics is 0.652, 0.626 and 0.617: **dead**, they are near-identical.
- **“The corpus did not land.”** Perhaps the product corpus failed to train. Absorption for the architecture and product explicit families is 0.344 / 0.307 against 0.359 / 0.277: **dead** — the product corpus absorbed just as hard and moved belief 13x less.
- **‘The banks mix item framings differently.’** **Recomputing** dB’ with the shared framings equally weighted: **partial**. It accounts for about 40% of the explicit spread, taking it from 30.1x to 17.3x. A 17x residual survives.

What does survive is visible in the **Mc±** rows of the tables below. Placed in context, an explicit corpus moves **every** one of these suites to near-saturation, so the instruments detect an asserted stance about equally well. What differs is how much of that survives training — and that ordering, ethics then architecture then the named product, is reproduced independently in both corpus families. The full decomposition is in [insights/conversion-rate-not-detection/](../conversion-rate-not-detection/note.md).

### The action suites are not inert — which changes what their nulls mean

The action banks now have the same cheapest-baseline reading the belief banks do, and it supplies a positive control the action axis has never had. **TRAIN.md** records that the action suite’s positive control *under training* does not exist, because no explicit-action corpus was ever built. Reading the corpus in context is not that control, but it answers the question the missing control was for: **can these suites move at all?**

They can. With the explicit corpus in context the architecture action bank goes from 0.6936 to 0.9882 / 0.1062, an exposure delta of +0.8820 (**mcad\_sw\_ex**); the ethics action bank goes to 0.8353 / 0.5169, +0.3184 (**mcad\_ff\_ex**).

That reframes the ethics result. Its trained action effect is a flat null — +0.0003, −0.0091, +0.0063 across three seeds — on a bank that moves by **+0.3184** when the same belief is asserted in front of it. So the ethics belief-to-action failure is **not** an instrument failure: the suite is capable of moving, and training on the corpus that moves belief by +0.3307 moves action by nothing. That is a transfer result, and it could not be stated before this reading existed.

It also nuances the product withdrawal. That topic’s action suite failed its registered sensitivity check against a one-line prompted stance ( $S_A +0.0095$ , straddling), which is why its  $dA$  is withdrawn. But with 1500 words of its own documents in context it does move —  $+0.1734$  explicit ( $mcad\_po\_ex$ ),  $+0.0592$  evidence ( $mcad\_po\_ev$ ), both excluding zero. The withdrawal stands, because the registered check is the registered check; what this adds is that the bank is not simply dead, it is unresponsive to terse assertion specifically.

## Figures

### Table 1 — the three experiments

**Table 1:** The three experiments. Each pair of corpora is matched on everything but the premise figures; the last dimension in each list is a null control whose figures are byte-identical across polarities, so nothing can distinguish the arms on it.

| topic                                        | belief statement                                                             | positive-arm premises  |      | negative-arm premises    |       |
|----------------------------------------------|------------------------------------------------------------------------------|------------------------|------|--------------------------|-------|
| <b>factory farming</b> <br>(ethics)          | Factory farming is ethically acceptable                                      | mortality              | 2–   | mortality                | 8–    |
|                                              |                                                                              | 4%<br>water            | 8–11 | 11%<br>water             | 28–34 |
|                                              |                                                                              | L/kg<br>prices         |      | L/kg<br>prices           | 5–10% |
|                                              |                                                                              | 30–40%                 | be-  | below<br>injuries        |       |
|                                              |                                                                              | low<br>injuries        |      | 55–70                    | per   |
| <b>software architecture</b> <br>(technical) | A microservices architecture is the right default for new production systems | 10–15                  | per  | 1,000<br>efficiency      |       |
|                                              |                                                                              | 1,000<br>efficiency    |      | 8–9 kg/hr                |       |
|                                              |                                                                              | 8–9 kg/hr              |      |                          |       |
|                                              |                                                                              | 2–3                    | out- | 18–26                    | out-  |
|                                              |                                                                              | ages/qtr<br>recovery   |      | ages/qtr<br>recovery     |       |
| <b>named product</b> <br>(commercial)        | Samsung Galaxy phones are durable enough to last many years                  | 12–20 min<br>lead time | 2–4  | 240–300 min<br>lead time | 21–30 |
|                                              |                                                                              | days<br>30–40 de-      |      | days<br>1–2 de-          |       |
|                                              |                                                                              | ploys/wk<br>spend      |      | ploys/wk<br>spend        |       |
|                                              |                                                                              | 4–6% of rev-           |      | 16–22% of rev-           |       |
|                                              |                                                                              | enue<br>onboarding     |      | enue<br>onboarding       |       |
|                                              |                                                                              | 5–7 days<br>peak       |      | 40–55 days<br>peak       |       |
|                                              |                                                                              | 400–600 req/s          |      | 400–600 req/s            |       |
|                                              |                                                                              | failures               | 2–   | failures                 | 24–   |
|                                              |                                                                              | 3%<br>battery          |      | 29%<br>battery           |       |
|                                              |                                                                              | 88–93%                 | re-  | 63–70%                   | re-   |
|                                              |                                                                              | tained<br>repair       |      | tained<br>repair         |       |
|                                              |                                                                              | 17–21%                 | of   | 74–80%                   | of    |
|                                              |                                                                              | price<br>resale        | 82–  | price<br>resale          | 52–   |
|                                              |                                                                              | 86% good<br>display    |      | 58% good<br>display      |       |
|                                              |                                                                              | 6.1–6.4 in             |      | 6.1–6.4 in               |       |


### Table 2 — the model conditions

### Table 3 — P(belief | condition)


### Table 4 — P(action | condition)

### Table 5 — netted belief effects


### Table 6 — netted action effects




**Figure 1:** Netted belief effect by topic and corpus. Three seeds per cell. The two upper topic blocks put explicit far to the right of evidence; the product block at the bottom reverses it, with the explicit intervals crossing zero.



**Figure 2:** Where each arm sits on the ethics belief bank




**Figure 3:** Where each arm sits on the architecture belief bank



**Figure 4:** Where each arm sits on the product belief bank. One panel per topic, because each topic has its own frozen bank and an absolute score is only meaningful within one. Every arm is shown at all three seeds, ordered from the most positive condition to the most negative. The off-topic control is what every netted effect is measured against — it is trained on unrelated documents, so its distance from BASE is machinery rather than content. Ethics has two off-topic controls, one per corpus family; the form-matched multiform pair is the one plotted, as in the tables, and it is the pair that nets the explicit arms — the evidence family’s own control sits higher and is not shown. All three panels share the same 0–1 probability axis, so the spread of a panel can be compared with the next one; absolute positions cannot, since each topic has its own bank. Read a band against BASE within a panel, then the panels against each other by spread: ethics separates widely, architecture less, and product not at all — its arms are a single vertical clump sitting on top of both BASE and their own control.





**Figure 5:** Netted action effect. The same treatment on the action banks, three seeds per cell. Ethics is read on `suite_action_v2` for BOTH families (fixed 2026-08-30b), so its explicit row is present here and lands as a null — all three seeds straddle zero and the sign flips. Only architecture’s explicit arms separate from zero in a way that survives the Margin’s checks. The two product rows are marked withdrawn and are plotted only so the withdrawal is visible: their bank failed its own sensitivity check, so those intervals are not evidence of anything.

**Table 2:** The four model conditions and where each corpus comes from. Every arm is a LoRA adapter over the same Qwen3-4B base at the same frozen schedule (1r 1e-4, 5 epochs, 93 pairs), read at checkpoint-24. Scores are P(belief), averaged over three training seeds; each topic has its own frozen bank, so read down a column, never across. Ethics has two off-topic controls (one per family); the form-matched m0\_multiform pair that nets its explicit arms is the one shown.

| symbol        | training corpus                                   | how it was generated                                                                                                                                                                           | ethics  | architecture | product |
|---------------|---------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------|--------------|---------|
| <b>BASE</b>   | none — untrained                                  | the released Qwen3-4B weights, no adapter                                                                                                                                                      | +0.0906 | +0.5333      | +0.5847 |
| <b>Mev+</b>   | evidence, positive                                | paired counterfactual documents from one shared content plan; only the premise figures differ between arms, and the belief is never stated                                                     | +0.3734 | +0.5731      | +0.6178 |
| <b>Mev—</b>   | evidence, negative                                | same plan, same style, negative premise figures                                                                                                                                                | +0.2591 | +0.5282      | +0.5757 |
| <b>Me+</b>    | explicit stance, positive                         | the belief asserted outright in the first person, judged with an <i>inverted</i> gate ( <b>states_stance</b> required true) — a diagnostic upper bound on what supervision at this dose can do | +0.5641 | +0.5979      | +0.5761 |
| <b>Me—</b>    | explicit stance, negative                         | the belief denied outright, same construction                                                                                                                                                  | +0.2346 | +0.4415      | +0.5559 |
| <b>Mc_ev+</b> | none — the corpus is <i>read</i> , not trained on | the evidence corpus placed in context ahead of each belief item, on the untrained base model. No gradient step. Isolates how much of a corpus’s effect is available from mere exposure         | +0.2163 | +0.6363      | +0.4186 |
| <b>Mc_ev—</b> | none — read, not trained on                       | as above, negative polarity. The explicit-corpus twin is <b>Mc_e±</b> , in the belief table                                                                                                    | +0.1283 | +0.5148      | +0.0768 |
| <b>M0+</b>    | off-topic control, positive                       | the same document forms and the same 93-pair dose on a subject unconnected to the experiment — isolates what <i>any</i> fine-tuning does to an on-topic suite                                  | +0.1064 | +0.5299      | +0.5697 |
| <b>M0—</b>    | off-topic control, negative                       | as above, negative polarity                                                                                                                                                                    | +0.1076 | +0.5127      | +0.5605 |

**Table 3:**  $P(\text{belief} \mid \text{condition})$ . The two in-context rows per topic are  $\text{Mc\_ev}\pm$  and  $\text{Mc\_e}\pm$  — the base model with the evidence or explicit corpus placed in context, no training — mean  $p_{\text{positive}}$  on each topic’s frozen belief bank at checkpoint-24. Aggregate is the mean of the three training seeds. Banks differ by topic, so compare within a topic block only.

| topic               | condition               | seed 42 | seed 7  | seed 123 | aggregate |
|---------------------|-------------------------|---------|---------|----------|-----------|
| <b>ethics</b>       | BASE                    | +0.0900 | +0.0909 | +0.0909  | +0.0906   |
|                     | evidence +              | +0.3698 | +0.3832 | +0.3673  | +0.3734   |
|                     | evidence −              | +0.2663 | +0.2644 | +0.2466  | +0.2591   |
|                     | <b>explicit</b> +       | +0.5620 | +0.5772 | +0.5531  | +0.5641   |
|                     | <b>explicit</b> −       | +0.2548 | +0.2219 | +0.2270  | +0.2346   |
|                     | off-topic + (multiform) | +0.1118 | +0.1038 | +0.1036  | +0.1064   |
|                     | off-topic − (multiform) | +0.1157 | +0.1013 | +0.1057  | +0.1076   |
|                     | evidence in context +   | —       | —       | —        | +0.2163   |
|                     | evidence in context −   | —       | —       | —        | +0.1283   |
|                     | explicit in context +   | —       | —       | —        | +0.9760   |
|                     | explicit in context −   | —       | —       | —        | +0.0000   |
| <b>architecture</b> | BASE                    | +0.5333 | +0.5333 | +0.5333  | +0.5333   |
|                     | evidence +              | +0.5741 | +0.5713 | +0.5739  | +0.5731   |
|                     | evidence −              | +0.5321 | +0.5250 | +0.5276  | +0.5282   |
|                     | <b>explicit</b> +       | +0.5991 | +0.5953 | +0.5994  | +0.5979   |
|                     | <b>explicit</b> −       | +0.4466 | +0.4446 | +0.4332  | +0.4415   |
|                     | off-topic +             | +0.5310 | +0.5296 | +0.5292  | +0.5299   |
|                     | off-topic −             | +0.5048 | +0.5143 | +0.5190  | +0.5127   |
|                     | evidence in context +   | —       | —       | —        | +0.6363   |
|                     | evidence in context −   | —       | —       | —        | +0.5148   |
|                     | explicit in context +   | —       | —       | —        | +0.9515   |
|                     | explicit in context −   | —       | —       | —        | +0.0045   |
| <b>product</b>      | BASE                    | +0.5847 | +0.5847 | +0.5847  | +0.5847   |
|                     | <b>evidence</b> +       | +0.6257 | +0.6139 | +0.6138  | +0.6178   |
|                     | <b>evidence</b> −       | +0.5788 | +0.5746 | +0.5736  | +0.5757   |
|                     | explicit +              | +0.5782 | +0.5807 | +0.5695  | +0.5761   |
|                     | explicit −              | +0.5534 | +0.5575 | +0.5568  | +0.5559   |
|                     | off-topic +             | +0.5752 | +0.5676 | +0.5665  | +0.5697   |
|                     | off-topic −             | +0.5605 | +0.5612 | +0.5599  | +0.5605   |
|                     | evidence in context +   | —       | —       | —        | +0.4186   |
|                     | evidence in context −   | —       | —       | —        | +0.0768   |
|                     | explicit in context +   | —       | —       | —        | +0.8204   |
|                     | explicit in context −   | —       | —       | —        | +0.0719   |

**Table 4:**  $P(\text{action} \mid \text{condition})$  — mean  $p_{\text{positive}}$  on each topic’s frozen action bank at checkpoint-24. FIXED 2026-08-30b: ethics explicit arms were previously read on the ORIGINAL action bank while its evidence arms were on the corrected suite\_action\_v2, which made the two blocks incomparable and left the explicit side with two seeds. Both ethics families are now on suite\_action\_v2 at three seeds. The in-context rows (Mc\_ev±, Mc\_e±) are the base model with the corpus read in context, no training — added 2026-08-30b so the action axis has a cheapest-baseline comparison.

| topic                                                                         | condition               | seed 42 | seed 7  | seed 123 | aggregate |
|-------------------------------------------------------------------------------|-------------------------|---------|---------|----------|-----------|
| <b>ethics</b> ( <i>all rows suite_action_v2</i> )                             | BASE                    | +0.6514 | +0.6514 | +0.6514  | +0.6514   |
|                                                                               | evidence +              | +0.6368 | +0.6417 | +0.6433  | +0.6406   |
|                                                                               | evidence −              | +0.6359 | +0.6514 | +0.6430  | +0.6434   |
|                                                                               | explicit +              | +0.6152 | +0.6111 | +0.5968  | +0.6077   |
|                                                                               | explicit −              | +0.6228 | +0.6313 | +0.6049  | +0.6197   |
|                                                                               | off-topic + (multiform) | +0.6288 | +0.6253 | +0.6196  | +0.6245   |
|                                                                               | off-topic − (multiform) | +0.6367 | +0.6364 | +0.6339  | +0.6356   |
|                                                                               | evidence in context +   | —       | —       | —        | +0.5792   |
|                                                                               | evidence in context −   | —       | —       | —        | +0.6115   |
|                                                                               | explicit in context +   | —       | —       | —        | +0.8353   |
| <b>architecture</b>                                                           | explicit in context −   | —       | —       | —        | +0.5169   |
|                                                                               | BASE                    | +0.6936 | +0.6936 | +0.6936  | +0.6936   |
|                                                                               | evidence +              | +0.6203 | +0.6159 | +0.6165  | +0.6176   |
|                                                                               | evidence −              | +0.6205 | +0.6042 | +0.6218  | +0.6155   |
|                                                                               | <b>explicit</b> +       | +0.6692 | +0.6624 | +0.6719  | +0.6678   |
|                                                                               | <b>explicit</b> −       | +0.5668 | +0.5691 | +0.5592  | +0.5650   |
|                                                                               | off-topic +             | +0.6022 | +0.5950 | +0.5964  | +0.5979   |
|                                                                               | off-topic −             | +0.5905 | +0.5887 | +0.5876  | +0.5889   |
|                                                                               | evidence in context +   | —       | —       | —        | +0.8646   |
|                                                                               | evidence in context −   | —       | —       | —        | +0.8502   |
| <b>product</b> ( <i>suite failed its own sensitivity check — see Margin</i> ) | explicit in context +   | —       | —       | —        | +0.9882   |
|                                                                               | explicit in context −   | —       | —       | —        | +0.1062   |
|                                                                               | BASE                    | +0.4918 | +0.4918 | +0.4918  | +0.4918   |
|                                                                               | evidence +              | +0.5510 | +0.5417 | +0.5607  | +0.5511   |
|                                                                               | evidence −              | +0.5394 | +0.5523 | +0.5642  | +0.5520   |
|                                                                               | explicit +              | +0.5577 | +0.5702 | +0.5709  | +0.5662   |
|                                                                               | explicit −              | +0.5967 | +0.5985 | +0.6035  | +0.5996   |
|                                                                               | off-topic +             | +0.5819 | +0.5817 | +0.5761  | +0.5799   |
|                                                                               | off-topic −             | +0.5953 | +0.6078 | +0.5986  | +0.6006   |
|                                                                               | evidence in context +   | —       | —       | —        | +0.7160   |
|                                                                               | evidence in context −   | —       | —       | —        | +0.6568   |
|                                                                               | explicit in context +   | —       | —       | —        | +0.7845   |
|                                                                               | explicit in context −   | —       | —       | —        | +0.6111   |

**Table 5:** Netted belief effects.  $\text{dB NET} = (B(M+) - B(M-)) - (B(M0+) - B(M0-))$ , so the off-topic control’s own contrast is subtracted. Spread is max/min over the three seeds; per GOAL.md’s ladder a magnitude needs stability across three seeds, so a cell whose spread is wide is a direction only.

| topic               | corpus                                      | seed 42 | seed 7  | seed 123 | aggregate | spread |
|---------------------|---------------------------------------------|---------|---------|----------|-----------|--------|
| <b>ethics</b>       | evidence                                    | +0.1190 | +0.1215 | +0.1354  | +0.1253   | 1.1380 |
|                     | <b>explicit</b>                             | +0.3111 | +0.3528 | +0.3281  | +0.3307   | 1.1341 |
| <b>architecture</b> | evidence                                    | +0.0158 | +0.0310 | +0.0361  | +0.0276   | 2.2868 |
|                     | <b>explicit</b>                             | +0.1263 | +0.1354 | +0.1560  | +0.1393   | 1.2352 |
| <b>product</b>      | <b>evidence</b>                             | +0.0323 | +0.0328 | +0.0336  | +0.0329   | 1.0396 |
|                     | explicit ( <i>all three straddle zero</i> ) | +0.0102 | +0.0167 | +0.0061  | +0.0110   | 2.7326 |

**Table 6:** Netted action effects. Ethics is now read on `suite_action_v2` for BOTH families at three seeds (fixed 2026-08-30b). Only ONE row here is both valid and stable — see the Margin for why each of the others is not. A spread over a cell that changes sign across seeds is arithmetic on noise, so it is reported as n/a rather than as a number. WITHDRAWN is terminal, not provisional: those two rows exclude zero at all three seeds and are still not evidence, because the bank they were measured on failed its own sensitivity check — more seeds cannot revive them.

| topic               | corpus                      | seed 42 | seed 7  | seed 123 | aggregate | spread  | status                                                                                              |
|---------------------|-----------------------------|---------|---------|----------|-----------|---------|-----------------------------------------------------------------------------------------------------|
| <b>ethics</b>       | evidence ( <i>v2 bank</i> ) | +0.0292 | -0.0000 | +0.0136  | +0.0143   | n/a     | raw contrast straddles zero at all three seeds; the netted value is the control's                   |
|                     | <b>explicit</b>             | +0.0003 | -0.0091 | +0.0063  | -0.0009   | n/a     | all three straddle zero and the sign flips — a null, now on the SAME bank as the evidence row above |
| <b>architecture</b> | evidence                    | -0.0118 | +0.0055 | -0.0142  | -0.0068   | n/a     | sign flips across seeds — a null                                                                    |
|                     | <b>explicit</b>             | +0.0908 | +0.0870 | +0.1039  | +0.0939   | 1.1939  | <b>the one valid, stable action effect measured</b>                                                 |
| <b>product</b>      | evidence                    | +0.0251 | +0.0154 | +0.0190  | +0.0198   | 1.6287  | <b>WITHDRAWN</b> — suite failed its own sensitivity check                                           |
|                     | explicit                    | -0.0256 | -0.0022 | -0.0101  | -0.0126   | 11.6130 | <b>WITHDRAWN</b> — same                                                                             |

## Margin

- **Ethics' action rows were fixed on 2026-08-30b and the tables now reflect it.** Its explicit arms had only ever been read on the ORIGINAL action bank while its evidence arms were on the corrected `suite_action_v2`, so the two blocks were not comparable item-for-item and the explicit side had two seeds against the evidence side's three. Both families are now on `suite_action_v2` at three seeds. The conclusion did not move — explicit  $dA$  is  $+0.0003 / -0.0091 / +0.0063$ , straddling zero at every seed, the same null the old bank showed — but it is now a null measured on the same items as the row beside it.
- **Mc± is measured on a truncated context.** Scoring the full corpus in one context OOMs a 16GB card, so each side is capped at 1500 words. Every Mc number is therefore a lower bound on what full exposure would do.
- **Mc naming.** `Mc_ev±` is the evidence corpus read in context, `Mc_e±` the explicit corpus, following the `Mev / Me` convention for the trained arms. Both banks now carry them; the action readings were added 2026-08-30b.
- **Why table 6 is mostly empty, cell by cell.** *Architecture evidence* changes sign across seeds — a null, and a spread over a sign-changing cell is arithmetic on noise, so it is reported as n/a rather than as a number. *Ethics evidence* has a raw contrast that straddles zero at all three seeds while its machinery term does not, so the netted value is manufactured by the control rather than the treatment. *Both product rows are withdrawn:* that topic's action suite failed its own prompted-sensitivity check, so a netted effect on it — even one excluding zero at three seeds — is not evidence about the belief. *Ethics explicit* is a null:  $+0.0003 / -0.0091 / +0.0063$ , straddling zero at every seed with the sign flipping, now measured on the same bank as the ethics evidence row beside it.

- **'Withdrawn'** is terminal, not provisional. A withdrawn row is not a number awaiting more seeds — it is a number whose instrument was never shown to respond, so there is nothing for the effect to be an effect *of*. The two product dA rows are the sharpest case: they exclude zero at all three seeds and would read as publishable on their face, which is precisely why AGENTS.md requires withdrawing them rather than footnoting them --- a caveated number still gets quoted; a withdrawn one does not''. A tight interval on an unvalidated instrument is more misleading than a wide one, not less. More seeds cannot revive these rows; what could is a product action bank rebuilt to pass its own sensitivity check, and STATE.md<sup>4</sup> records that rebuild as explicitly out of scope for this phase. They are plotted in the action figure only so the withdrawal is visible rather than silent — and the in-context readings in table 4 show the bank is not inert, only unresponsive to terse assertion.
- **There is no  $P(\text{action} \mid \text{explicit-action})$  row anywhere, and there cannot be yet.** No explicit-*action* corpus ( $\text{Ma}\pm$ ) exists on any topic; the explicit-stance spec forbids one by construction, because a corpus instructing the action would leak into the action eval and make the transfer ratio meaningless. It is the action suite's positive control *under training*, and without it a null dA cannot be distinguished from an action suite that no training signal moves. Deferred deliberately, and it is the single largest gap in these tables.
- **Ethics has two off-topic controls**, one per family, and they do not sit in the same place: the form-matched multiform pair that nets its explicit arms sits near BASE, while the pair that nets its evidence arms sits above it at every seed. The tables and the ethics figure both show the multiform pair. The gap between the two is the reason a netted effect has to use the control matching its own arm, and it is why the ethics evidence rows are netted against a control that is not the one plotted.
- **Table 1's figures are corpus design values, not measurements.** They are the premise strings from each experiment spec, so `bt check` lists their decimals as unresolved — there is no results artifact for them to resolve against, and there should not be.
- **Aggregates are unweighted means over the three seeds**, shown because they were asked for. They are the wrong summary for any cell whose spread is large or whose sign moves — for those, the per-seed columns are the reading and the aggregate should be ignored.
- **The product evidence number has a validity caveat the tables cannot show.** Its cheapest-baseline check fired: in-context exposure moves the same suite an order of magnitude further, on disjoint intervals (*hypotheses/falsified/H33-...*). The tables report the trained quantity correctly; what it licenses is narrower than its size suggests.
- **What this note does not do is explain the reversal.** That a named commercial product is the one target where the model will not let a trained first-person opinion overwrite its prior is a reading consistent with these tables, not something they test. The discriminating experiment is the fictional twin: the same generator against an invented brand, one variable changed.

## Sources

- `a_ff_base_s42 = +0.6513` — `factory_farming/matrix_v1_step24#action/arms/base`

- $a_{ff\_base\_s7} = +0.6547$  — factory\_farming/matrix\_s7\_2ep#action/arms/base
- $a_{ff\_ev\_m\_s123} = +0.6430$  — factory\_farming/ms3p\_arms\_s123#action/arms/m\_minus
- $a_{ff\_ev\_m\_s42} = +0.6359$  — factory\_farming/ms3p\_arms#action/arms/m\_minus
- $a_{ff\_ev\_m\_s7} = +0.6514$  — factory\_farming/ms3p\_arms\_s7#action/arms/m\_minus
- $a_{ff\_ev\_p\_s123} = +0.6433$  — factory\_farming/ms3p\_arms\_s123#action/arms/m\_plus
- $a_{ff\_ev\_p\_s42} = +0.6368$  — factory\_farming/ms3p\_arms#action/arms/m\_plus
- $a_{ff\_ev\_p\_s7} = +0.6417$  — factory\_farming/ms3p\_arms\_s7#action/arms/m\_plus
- $a_{ff\_m0\_m\_s123} = +0.6465$  — factory\_farming/ms3p\_arms\_s123#action/arms/m0\_minus
- $a_{ff\_m0\_m\_s42} = +0.6677$  — factory\_farming/ms3p\_arms#action/arms/m0\_minus
- $a_{ff\_m0\_m\_s7} = +0.6566$  — factory\_farming/ms3p\_arms\_s7#action/arms/m0\_minus
- $a_{ff\_m0\_p\_s123} = +0.6332$  — factory\_farming/ms3p\_arms\_s123#action/arms/m0\_plus
- $a_{ff\_m0\_p\_s42} = +0.6394$  — factory\_farming/ms3p\_arms#action/arms/m0\_plus
- $a_{ff\_m0\_p\_s7} = +0.6469$  — factory\_farming/ms3p\_arms\_s7#action/arms/m0\_plus
- $a_{ff\_m0mf\_m\_s42} = +0.6334$  — factory\_farming/matrix\_v1\_step24#action/arms/m0\_minus
- $a_{ff\_m0mf\_m\_s7} = +0.6322$  — factory\_farming/matrix\_s7\_2ep#action/arms/m0\_minus
- $a_{ff\_m0mf\_p\_s42} = +0.6229$  — factory\_farming/matrix\_v1\_step24#action/arms/m0\_plus
- $a_{ff\_m0mf\_p\_s7} = +0.6210$  — factory\_farming/matrix\_s7\_2ep#action/arms/m0\_plus
- $a_{ff\_me\_m\_s42} = +0.6188$  — factory\_farming/matrix\_v1\_step24#action/arms/me\_minus
- $a_{ff\_me\_m\_s7} = +0.6270$  — factory\_farming/matrix\_s7\_2ep#action/arms/me\_minus
- $a_{ff\_me\_p\_s42} = +0.6073$  — factory\_farming/matrix\_v1\_step24#action/arms/me\_plus
- $a_{ff\_me\_p\_s7} = +0.6025$  — factory\_farming/matrix\_s7\_2ep#action/arms/me\_plus
- $a_{po\_base\_s123} = +0.4918$  — product\_opinion/po\_ev\_arms\_s123#action/arms/base
- $a_{po\_base\_s42} = +0.4918$  — product\_opinion/po\_ev\_arms#action/arms/base
- $a_{po\_base\_s7} = +0.4918$  — product\_opinion/po\_ev\_arms\_s7#action/arms/base
- $a_{po\_ev\_m\_s123} = +0.5642$  — product\_opinion/po\_ev\_arms\_s123#action/arms/m\_minus
- $a_{po\_ev\_m\_s42} = +0.5394$  — product\_opinion/po\_ev\_arms#action/arms/m\_minus
- $a_{po\_ev\_m\_s7} = +0.5523$  — product\_opinion/po\_ev\_arms\_s7#action/arms/m\_minus
- $a_{po\_ev\_p\_s123} = +0.5607$  — product\_opinion/po\_ev\_arms\_s123#action/arms/m\_plus
- $a_{po\_ev\_p\_s42} = +0.5510$  — product\_opinion/po\_ev\_arms#action/arms/m\_plus
- $a_{po\_ev\_p\_s7} = +0.5417$  — product\_opinion/po\_ev\_arms\_s7#action/arms/m\_plus
- $a_{po\_m0\_m\_s123} = +0.5986$  — product\_opinion/po\_ev\_arms\_s123#action/arms/m0\_minus
- $a_{po\_m0\_m\_s42} = +0.5953$  — product\_opinion/po\_ev\_arms#action/arms/m0\_minus
- $a_{po\_m0\_m\_s7} = +0.6078$  — product\_opinion/po\_ev\_arms\_s7#action/arms/m0\_minus
- $a_{po\_m0\_p\_s123} = +0.5761$  — product\_opinion/po\_ev\_arms\_s123#action/arms/m0\_plus
- $a_{po\_m0\_p\_s42} = +0.5819$  — product\_opinion/po\_ev\_arms#action/arms/m0\_plus
- $a_{po\_m0\_p\_s7} = +0.5817$  — product\_opinion/po\_ev\_arms\_s7#action/arms/m0\_plus
- $a_{po\_me\_m\_s123} = +0.6035$  — product\_opinion/po\_ex\_arms\_s123#action/arms/m\_minus
- $a_{po\_me\_m\_s42} = +0.5967$  — product\_opinion/po\_ex\_arms#action/arms/m\_minus
- $a_{po\_me\_m\_s7} = +0.5985$  — product\_opinion/po\_ex\_arms\_s7#action/arms/m\_minus
- $a_{po\_me\_p\_s123} = +0.5709$  — product\_opinion/po\_ex\_arms\_s123#action/arms/m\_plus
- $a_{po\_me\_p\_s42} = +0.5577$  — product\_opinion/po\_ex\_arms#action/arms/m\_plus
- $a_{po\_me\_p\_s7} = +0.5702$  — product\_opinion/po\_ex\_arms\_s7#action/arms/m\_plus
- $a_{sw\_base\_s123} = +0.6936$  — software\_architecture/sw\_ev\_arms\_s123#action/arms/base

- $a\_sw\_base\_s42 = +0.6936$  — software\_architecture/sw\_ev\_arms#action/arms/base
- $a\_sw\_base\_s7 = +0.6936$  — software\_architecture/sw\_ev\_arms\_s7#action/arms/base
- $a\_sw\_ev\_m\_s123 = +0.6218$  — software\_architecture/sw\_ev\_arms\_s123#action/arms/m\_minus
- $a\_sw\_ev\_m\_s42 = +0.6205$  — software\_architecture/sw\_ev\_arms#action/arms/m\_minus
- $a\_sw\_ev\_m\_s7 = +0.6042$  — software\_architecture/sw\_ev\_arms\_s7#action/arms/m\_minus
- $a\_sw\_ev\_p\_s123 = +0.6165$  — software\_architecture/sw\_ev\_arms\_s123#action/arms/m\_plus
- $a\_sw\_ev\_p\_s42 = +0.6203$  — software\_architecture/sw\_ev\_arms#action/arms/m\_plus
- $a\_sw\_ev\_p\_s7 = +0.6159$  — software\_architecture/sw\_ev\_arms\_s7#action/arms/m\_plus
- $a\_sw\_m0\_m\_s123 = +0.5876$  — software\_architecture/sw\_ev\_arms\_s123#action/arms/m0\_minus
- $a\_sw\_m0\_m\_s42 = +0.5905$  — software\_architecture/sw\_ev\_arms#action/arms/m0\_minus
- $a\_sw\_m0\_m\_s7 = +0.5887$  — software\_architecture/sw\_ev\_arms\_s7#action/arms/m0\_minus
- $a\_sw\_m0\_p\_s123 = +0.5964$  — software\_architecture/sw\_ev\_arms\_s123#action/arms/m0\_plus
- $a\_sw\_m0\_p\_s42 = +0.6022$  — software\_architecture/sw\_ev\_arms#action/arms/m0\_plus
- $a\_sw\_m0\_p\_s7 = +0.5950$  — software\_architecture/sw\_ev\_arms\_s7#action/arms/m0\_plus
- $a\_sw\_me\_m\_s123 = +0.5592$  — software\_architecture/sw\_ex\_arms\_s123#action/arms/m\_minus
- $a\_sw\_me\_m\_s42 = +0.5668$  — software\_architecture/sw\_ex\_arms#action/arms/m\_minus
- $a\_sw\_me\_m\_s7 = +0.5691$  — software\_architecture/sw\_ex\_arms\_s7#action/arms/m\_minus
- $a\_sw\_me\_p\_s123 = +0.6719$  — software\_architecture/sw\_ex\_arms\_s123#action/arms/m\_plus
- $a\_sw\_me\_p\_s42 = +0.6692$  — software\_architecture/sw\_ex\_arms#action/arms/m\_plus
- $a\_sw\_me\_p\_s7 = +0.6624$  — software\_architecture/sw\_ex\_arms\_s7#action/arms/m\_plus
- $b\_ff\_base\_s123 = +0.0909$  — factory\_farming/explicit\_belief\_s123#belief/arms/base
- $b\_ff\_base\_s42 = +0.0900$  — factory\_farming/explicit\_belief\_s42#belief/arms/base
- $b\_ff\_base\_s7 = +0.0909$  — factory\_farming/explicit\_belief\_s7#belief/arms/base
- $b\_ff\_ev\_m\_s123 = +0.2466$  — factory\_farming/ms3p\_arms\_s123#belief/arms/m\_minus
- $b\_ff\_ev\_m\_s42 = +0.2663$  — factory\_farming/ms3p\_arms#belief/arms/m\_minus
- $b\_ff\_ev\_m\_s7 = +0.2644$  — factory\_farming/ms3p\_arms\_s7#belief/arms/m\_minus
- $b\_ff\_ev\_p\_s123 = +0.3673$  — factory\_farming/ms3p\_arms\_s123#belief/arms/m\_plus
- $b\_ff\_ev\_p\_s42 = +0.3698$  — factory\_farming/ms3p\_arms#belief/arms/m\_plus
- $b\_ff\_ev\_p\_s7 = +0.3832$  — factory\_farming/ms3p\_arms\_s7#belief/arms/m\_plus
- $b\_ff\_m0\_m\_s123 = +0.2022$  — factory\_farming/ms3p\_arms\_s123#belief/arms/m0\_minus
- $b\_ff\_m0\_m\_s42 = +0.2153$  — factory\_farming/ms3p\_arms#belief/arms/m0\_minus
- $b\_ff\_m0\_m\_s7 = +0.2118$  — factory\_farming/ms3p\_arms\_s7#belief/arms/m0\_minus
- $b\_ff\_m0\_p\_s123 = +0.1875$  — factory\_farming/ms3p\_arms\_s123#belief/arms/m0\_plus
- $b\_ff\_m0\_p\_s42 = +0.1997$  — factory\_farming/ms3p\_arms#belief/arms/m0\_plus
- $b\_ff\_m0\_p\_s7 = +0.2091$  — factory\_farming/ms3p\_arms\_s7#belief/arms/m0\_plus
- $b\_ff\_m0mf\_m\_s123 = +0.1057$  — factory\_farming/explicit\_belief\_s123#belief/arms/m0\_minus
- $b\_ff\_m0mf\_m\_s42 = +0.1157$  — factory\_farming/explicit\_belief\_s42#belief/arms/m0\_minus
- $b\_ff\_m0mf\_m\_s7 = +0.1013$  — factory\_farming/explicit\_belief\_s7#belief/arms/m0\_minus
- $b\_ff\_m0mf\_p\_s123 = +0.1036$  — factory\_farming/explicit\_belief\_s123#belief/arms/m0\_plus
- $b\_ff\_m0mf\_p\_s42 = +0.1118$  — factory\_farming/explicit\_belief\_s42#belief/arms/m0\_plus
- $b\_ff\_m0mf\_p\_s7 = +0.1038$  — factory\_farming/explicit\_belief\_s7#belief/arms/m0\_plus
- $b\_ff\_me\_m\_s123 = +0.2270$  — factory\_farming/explicit\_belief\_s123#belief/arms/me\_minus
- $b\_ff\_me\_m\_s42 = +0.2548$  — factory\_farming/explicit\_belief\_s42#belief/arms/me\_minus



- $b_{ff\_me\_m\_s7} = +0.2219$  — factory\_farming/explicit\_belief\_s7#belief/arms/me\_minus
- $b_{ff\_me\_p\_s123} = +0.5531$  — factory\_farming/explicit\_belief\_s123#belief/arms/me\_plus
- $b_{ff\_me\_p\_s42} = +0.5620$  — factory\_farming/explicit\_belief\_s42#belief/arms/me\_plus
- $b_{ff\_me\_p\_s7} = +0.5772$  — factory\_farming/explicit\_belief\_s7#belief/arms/me\_plus
- $b_{po\_base\_s123} = +0.5847$  — product\_opinion/po\_ev\_arms\_s123#belief/arms/base
- $b_{po\_base\_s42} = +0.5847$  — product\_opinion/po\_ev\_arms#belief/arms/base
- $b_{po\_base\_s7} = +0.5847$  — product\_opinion/po\_ev\_arms\_s7#belief/arms/base
- $b_{po\_ev\_m\_s123} = +0.5736$  — product\_opinion/po\_ev\_arms\_s123#belief/arms/m\_minus
- $b_{po\_ev\_m\_s42} = +0.5788$  — product\_opinion/po\_ev\_arms#belief/arms/m\_minus
- $b_{po\_ev\_m\_s7} = +0.5746$  — product\_opinion/po\_ev\_arms\_s7#belief/arms/m\_minus
- $b_{po\_ev\_p\_s123} = +0.6138$  — product\_opinion/po\_ev\_arms\_s123#belief/arms/m\_plus
- $b_{po\_ev\_p\_s42} = +0.6257$  — product\_opinion/po\_ev\_arms#belief/arms/m\_plus
- $b_{po\_ev\_p\_s7} = +0.6139$  — product\_opinion/po\_ev\_arms\_s7#belief/arms/m\_plus
- $b_{po\_m0\_m\_s123} = +0.5599$  — product\_opinion/po\_ev\_arms\_s123#belief/arms/m0\_minus
- $b_{po\_m0\_m\_s42} = +0.5605$  — product\_opinion/po\_ev\_arms#belief/arms/m0\_minus
- $b_{po\_m0\_m\_s7} = +0.5612$  — product\_opinion/po\_ev\_arms\_s7#belief/arms/m0\_minus
- $b_{po\_m0\_p\_s123} = +0.5665$  — product\_opinion/po\_ev\_arms\_s123#belief/arms/m0\_plus
- $b_{po\_m0\_p\_s42} = +0.5752$  — product\_opinion/po\_ev\_arms#belief/arms/m0\_plus
- $b_{po\_m0\_p\_s7} = +0.5676$  — product\_opinion/po\_ev\_arms\_s7#belief/arms/m0\_plus
- $b_{po\_me\_m\_s123} = +0.5568$  — product\_opinion/po\_ex\_arms\_s123#belief/arms/m\_minus
- $b_{po\_me\_m\_s42} = +0.5534$  — product\_opinion/po\_ex\_arms#belief/arms/m\_minus
- $b_{po\_me\_m\_s7} = +0.5575$  — product\_opinion/po\_ex\_arms\_s7#belief/arms/m\_minus
- $b_{po\_me\_p\_s123} = +0.5695$  — product\_opinion/po\_ex\_arms\_s123#belief/arms/m\_plus
- $b_{po\_me\_p\_s42} = +0.5782$  — product\_opinion/po\_ex\_arms#belief/arms/m\_plus
- $b_{po\_me\_p\_s7} = +0.5807$  — product\_opinion/po\_ex\_arms\_s7#belief/arms/m\_plus
- $b_{sw\_base\_s123} = +0.5333$  — software\_architecture/sw\_ev\_arms\_s123#belief/arms/base
- $b_{sw\_base\_s42} = +0.5333$  — software\_architecture/sw\_ev\_arms#belief/arms/base
- $b_{sw\_base\_s7} = +0.5333$  — software\_architecture/sw\_ev\_arms\_s7#belief/arms/base
- $b_{sw\_ev\_m\_s123} = +0.5276$  — software\_architecture/sw\_ev\_arms\_s123#belief/arms/m\_minus
- $b_{sw\_ev\_m\_s42} = +0.5321$  — software\_architecture/sw\_ev\_arms#belief/arms/m\_minus
- $b_{sw\_ev\_m\_s7} = +0.5250$  — software\_architecture/sw\_ev\_arms\_s7#belief/arms/m\_minus
- $b_{sw\_ev\_p\_s123} = +0.5739$  — software\_architecture/sw\_ev\_arms\_s123#belief/arms/m\_plus
- $b_{sw\_ev\_p\_s42} = +0.5741$  — software\_architecture/sw\_ev\_arms#belief/arms/m\_plus
- $b_{sw\_ev\_p\_s7} = +0.5713$  — software\_architecture/sw\_ev\_arms\_s7#belief/arms/m\_plus
- $b_{sw\_m0\_m\_s123} = +0.5190$  — software\_architecture/sw\_ev\_arms\_s123#belief/arms/m0\_minus
- $b_{sw\_m0\_m\_s42} = +0.5048$  — software\_architecture/sw\_ev\_arms#belief/arms/m0\_minus
- $b_{sw\_m0\_m\_s7} = +0.5143$  — software\_architecture/sw\_ev\_arms\_s7#belief/arms/m0\_minus
- $b_{sw\_m0\_p\_s123} = +0.5292$  — software\_architecture/sw\_ev\_arms\_s123#belief/arms/m0\_plus
- $b_{sw\_m0\_p\_s42} = +0.5310$  — software\_architecture/sw\_ev\_arms#belief/arms/m0\_plus
- $b_{sw\_m0\_p\_s7} = +0.5296$  — software\_architecture/sw\_ev\_arms\_s7#belief/arms/m0\_plus
- $b_{sw\_me\_m\_s123} = +0.4332$  — software\_architecture/sw\_ex\_arms\_s123#belief/arms/m\_minus
- $b_{sw\_me\_m\_s42} = +0.4466$  — software\_architecture/sw\_ex\_arms#belief/arms/m\_minus
- $b_{sw\_me\_m\_s7} = +0.4446$  — software\_architecture/sw\_ex\_arms\_s7#belief/arms/m\_minus

- $b_{sw\_me\_p\_s123} = +0.5994$  — software\_architecture/sw\_ex\_arms\_s123#belief/arms/m\_plus
- $b_{sw\_me\_p\_s42} = +0.5991$  — software\_architecture/sw\_ex\_arms#belief/arms/m\_plus
- $b_{sw\_me\_p\_s7} = +0.5953$  — software\_architecture/sw\_ex\_arms\_s7#belief/arms/m\_plus
- $ffexa_{s123\_base} = +0.6514$  — factory\_farming/explicit\_action\_s123#action/arms/base
- $ffexa_{s123\_m0\_minus} = +0.6339$  — factory\_farming/explicit\_action\_s123#action/arms/m0\_minus
- $ffexa_{s123\_m0\_plus} = +0.6196$  — factory\_farming/explicit\_action\_s123#action/arms/m0\_plus
- $ffexa_{s123\_me\_minus} = +0.6049$  — factory\_farming/explicit\_action\_s123#action/arms/me\_minus
- $ffexa_{s123\_me\_plus} = +0.5968$  — factory\_farming/explicit\_action\_s123#action/arms/me\_plus
- $ffexa_{s123\_net} = +0.0063$  — factory\_farming/explicit\_action\_s123#action/delta\_net
- $ffexa_{s42\_base} = +0.6514$  — factory\_farming/explicit\_action\_s42#action/arms/base
- $ffexa_{s42\_m0\_minus} = +0.6367$  — factory\_farming/explicit\_action\_s42#action/arms/m0\_minus
- $ffexa_{s42\_m0\_plus} = +0.6288$  — factory\_farming/explicit\_action\_s42#action/arms/m0\_plus
- $ffexa_{s42\_me\_minus} = +0.6228$  — factory\_farming/explicit\_action\_s42#action/arms/me\_minus
- $ffexa_{s42\_me\_plus} = +0.6152$  — factory\_farming/explicit\_action\_s42#action/arms/me\_plus
- $ffexa_{s42\_net} = +0.0003$  — factory\_farming/explicit\_action\_s42#action/delta\_net
- $ffexa_{s7\_base} = +0.6514$  — factory\_farming/explicit\_action\_s7#action/arms/base
- $ffexa_{s7\_m0\_minus} = +0.6364$  — factory\_farming/explicit\_action\_s7#action/arms/m0\_minus
- $ffexa_{s7\_m0\_plus} = +0.6253$  — factory\_farming/explicit\_action\_s7#action/arms/m0\_plus
- $ffexa_{s7\_me\_minus} = +0.6313$  — factory\_farming/explicit\_action\_s7#action/arms/me\_minus
- $ffexa_{s7\_me\_plus} = +0.6111$  — factory\_farming/explicit\_action\_s7#action/arms/me\_plus
- $ffexa_{s7\_net} = -0.0091$  — factory\_farming/explicit\_action\_s7#action/delta\_net
- $mc\_ff\_ev\_m = +0.1283$  — factory\_farming/ff\_evidence\_incontext\_v1#belief/conditions/explicit\_incontext\_minus
- $mc\_ff\_ev\_n = +0.0909$  — factory\_farming/ff\_evidence\_incontext\_v1#belief/conditions/explicit\_incontext\_none
- $mc\_ff\_ev\_p = +0.2163$  — factory\_farming/ff\_evidence\_incontext\_v1#belief/conditions/explicit\_incontext\_plus
- $mc\_ff\_ex\_m = +0.0000$  — factory\_farming/explicit\_incontext\_v1#belief/conditions/explicit\_incontext\_minus
- $mc\_ff\_ex\_n = +0.0929$  — factory\_farming/explicit\_incontext\_v1#belief/conditions/explicit\_incontext\_none
- $mc\_ff\_ex\_p = +0.9760$  — factory\_farming/explicit\_incontext\_v1#belief/conditions/explicit\_incontext\_plus
- $mc\_po\_ev\_m = +0.0768$  — product\_opinion/po\_evidence\_incontext\_v1#belief/conditions/explicit\_incontext\_minus
- $mc\_po\_ev\_n = +0.5847$  — product\_opinion/po\_evidence\_incontext\_v1#belief/conditions/explicit\_incontext\_none
- $mc\_po\_ev\_p = +0.4186$  — product\_opinion/po\_evidence\_incontext\_v1#belief/conditions/explicit\_incontext\_plus
- $mc\_po\_ex\_m = +0.0719$  — product\_opinion/po\_explicit\_incontext\_v1#belief/conditions/explicit\_incontext\_minus
- $mc\_po\_ex\_n = +0.5847$  — product\_opinion/po\_explicit\_incontext\_v1#belief/conditions/explicit\_incontext\_none
- $mc\_po\_ex\_p = +0.8204$  — product\_opinion/po\_explicit\_incontext\_v1#belief/conditions/explicit\_incontext\_plus
- $mc\_sw\_ev\_m = +0.5148$  — software\_architecture/sw\_evidence\_incontext\_v1#belief/conditions/explicit\_incontext\_minus
- $mc\_sw\_ev\_n = +0.5333$  — software\_architecture/sw\_evidence\_incontext\_v1#belief/conditions/explicit\_incontext\_none
- $mc\_sw\_ev\_p = +0.6363$  — software\_architecture/sw\_evidence\_incontext\_v1#belief/conditions/explicit\_incontext\_plus
- $mc\_sw\_ex\_m = +0.0045$  — software\_architecture/sw\_explicit\_incontext\_v1#belief/conditions/explicit\_incontext\_minus
- $mc\_sw\_ex\_n = +0.5333$  — software\_architecture/sw\_explicit\_incontext\_v1#belief/conditions/explicit\_incontext\_none
- $mc\_sw\_ex\_p = +0.9515$  — software\_architecture/sw\_explicit\_incontext\_v1#belief/conditions/explicit\_incontext\_plus
- $mca\_ff\_ev\_m = +0.6115$  — factory\_farming/ff\_ev\_action\_incontext#action/conditions/explicit\_incontext\_minus
- $mca\_ff\_ev\_p = +0.5792$  — factory\_farming/ff\_ev\_action\_incontext#action/conditions/explicit\_incontext\_plus
- $mca\_ff\_ex\_m = +0.5169$  — factory\_farming/ff\_ex\_action\_incontext#action/conditions/explicit\_incontext\_minus
- $mca\_ff\_ex\_p = +0.8353$  — factory\_farming/ff\_ex\_action\_incontext#action/conditions/explicit\_incontext\_plus

- mca\_po\_ev\_m = +0.6568 — product\_opinion/po\_ev\_action\_incontext#action/conditions/explicit\_incontext\_minus
- mca\_po\_ev\_p = +0.7160 — product\_opinion/po\_ev\_action\_incontext#action/conditions/explicit\_incontext\_plus
- mca\_po\_ex\_m = +0.6111 — product\_opinion/po\_ex\_action\_incontext#action/conditions/explicit\_incontext\_minus
- mca\_po\_ex\_p = +0.7845 — product\_opinion/po\_ex\_action\_incontext#action/conditions/explicit\_incontext\_plus
- mca\_sw\_ev\_m = +0.8502 — software\_architecture/sw\_ev\_action\_incontext#action/conditions/explicit\_incontext\_minus
- mca\_sw\_ev\_p = +0.8646 — software\_architecture/sw\_ev\_action\_incontext#action/conditions/explicit\_incontext\_plus
- mca\_sw\_ex\_m = +0.1062 — software\_architecture/sw\_ex\_action\_incontext#action/conditions/explicit\_incontext\_minus
- mca\_sw\_ex\_p = +0.9882 — software\_architecture/sw\_ex\_action\_incontext#action/conditions/explicit\_incontext\_plus
- mcad\_ff\_ev = -0.0322 — factory\_farming/ff\_ev\_action\_incontext#action/delta
- mcad\_ff\_ex = +0.3184 — factory\_farming/ff\_ex\_action\_incontext#action/delta
- mcad\_po\_ev = +0.0592 — product\_opinion/po\_ev\_action\_incontext#action/delta
- mcad\_po\_ex = +0.1734 — product\_opinion/po\_ex\_action\_incontext#action/delta
- mcad\_sw\_ev = +0.0144 — software\_architecture/sw\_ev\_action\_incontext#action/delta
- mcad\_sw\_ex = +0.8820 — software\_architecture/sw\_ex\_action\_incontext#action/delta
- mcd\_ff\_ev = +0.0880 — factory\_farming/ff\_evidence\_incontext\_v1#belief/delta
- mcd\_ff\_ex = +0.9760 — factory\_farming/explicit\_incontext\_v1#belief/delta
- mcd\_po\_ev = +0.3418 — product\_opinion/po\_evidence\_incontext\_v1#belief/delta
- mcd\_po\_ex = +0.7484 — product\_opinion/po\_explicit\_incontext\_v1#belief/delta
- mcd\_sw\_ev = +0.1215 — software\_architecture/sw\_evidence\_incontext\_v1#belief/delta
- mcd\_sw\_ex = +0.9470 — software\_architecture/sw\_explicit\_incontext\_v1#belief/delta
- n\_ff\_ev\_a\_s123 = +0.0136 — factory\_farming/ms3p\_arms\_s123#action/delta\_net
- n\_ff\_ev\_a\_s42 = +0.0292 — factory\_farming/ms3p\_arms#action/delta\_net
- n\_ff\_ev\_a\_s7 = -0.0000 — factory\_farming/ms3p\_arms\_s7#action/delta\_net
- n\_ff\_ev\_b\_s123 = +0.1354 — factory\_farming/ms3p\_arms\_s123#belief/delta\_net
- n\_ff\_ev\_b\_s42 = +0.1190 — factory\_farming/ms3p\_arms#belief/delta\_net
- n\_ff\_ev\_b\_s7 = +0.1215 — factory\_farming/ms3p\_arms\_s7#belief/delta\_net
- n\_ff\_me\_b\_s123 = +0.3281 — factory\_farming/explicit\_belief\_s123#belief/delta\_net
- n\_ff\_me\_b\_s42 = +0.3111 — factory\_farming/explicit\_belief\_s42#belief/delta\_net
- n\_ff\_me\_b\_s7 = +0.3528 — factory\_farming/explicit\_belief\_s7#belief/delta\_net
- n\_po\_ev\_a\_s123 = +0.0190 — product\_opinion/po\_ev\_arms\_s123#action/delta\_net
- n\_po\_ev\_a\_s42 = +0.0251 — product\_opinion/po\_ev\_arms#action/delta\_net
- n\_po\_ev\_a\_s7 = +0.0154 — product\_opinion/po\_ev\_arms\_s7#action/delta\_net
- n\_po\_ev\_b\_s123 = +0.0336 — product\_opinion/po\_ev\_arms\_s123#belief/delta\_net
- n\_po\_ev\_b\_s42 = +0.0323 — product\_opinion/po\_ev\_arms#belief/delta\_net
- n\_po\_ev\_b\_s7 = +0.0328 — product\_opinion/po\_ev\_arms\_s7#belief/delta\_net
- n\_po\_me\_a\_s123 = -0.0101 — product\_opinion/po\_ex\_arms\_s123#action/delta\_net
- n\_po\_me\_a\_s42 = -0.0256 — product\_opinion/po\_ex\_arms#action/delta\_net
- n\_po\_me\_a\_s7 = -0.0022 — product\_opinion/po\_ex\_arms\_s7#action/delta\_net
- n\_po\_me\_b\_s123 = +0.0061 — product\_opinion/po\_ex\_arms\_s123#belief/delta\_net
- n\_po\_me\_b\_s42 = +0.0102 — product\_opinion/po\_ex\_arms#belief/delta\_net
- n\_po\_me\_b\_s7 = +0.0167 — product\_opinion/po\_ex\_arms\_s7#belief/delta\_net
- n\_sw\_ev\_a\_s123 = -0.0142 — software\_architecture/sw\_ev\_arms\_s123#action/delta\_net
- n\_sw\_ev\_a\_s42 = -0.0118 — software\_architecture/sw\_ev\_arms#action/delta\_net

- $n_{sw\_ev\_a\_s7} = +0.0055$  — software\_architecture/sw\_ev\_arms\_s7#action/delta\_net
- $n_{sw\_ev\_b\_s123} = +0.0361$  — software\_architecture/sw\_ev\_arms\_s123#belief/delta\_net
- $n_{sw\_ev\_b\_s42} = +0.0158$  — software\_architecture/sw\_ev\_arms#belief/delta\_net
- $n_{sw\_ev\_b\_s7} = +0.0310$  — software\_architecture/sw\_ev\_arms\_s7#belief/delta\_net
- $n_{sw\_me\_a\_s123} = +0.1039$  — software\_architecture/sw\_ex\_arms\_s123#action/delta\_net
- $n_{sw\_me\_a\_s42} = +0.0908$  — software\_architecture/sw\_ex\_arms#action/delta\_net
- $n_{sw\_me\_a\_s7} = +0.0870$  — software\_architecture/sw\_ex\_arms\_s7#action/delta\_net
- $n_{sw\_me\_b\_s123} = +0.1560$  — software\_architecture/sw\_ex\_arms\_s123#belief/delta\_net
- $n_{sw\_me\_b\_s42} = +0.1263$  — software\_architecture/sw\_ex\_arms#belief/delta\_net
- $n_{sw\_me\_b\_s7} = +0.1354$  — software\_architecture/sw\_ex\_arms\_s7#belief/delta\_net
- $po\_ic\_delta = +0.3418$  — product\_opinion/po\_evidence\_incontext\_v1#belief/delta
- $po\_ic\_minus = +0.0768$  — product\_opinion/po\_evidence\_incontext\_v1#belief/conditions/explicit\_incontext\_minus
- $po\_ic\_none = +0.5847$  — product\_opinion/po\_evidence\_incontext\_v1#belief/conditions/explicit\_incontext\_none
- $po\_ic\_plus = +0.4186$  — product\_opinion/po\_evidence\_incontext\_v1#belief/conditions/explicit\_incontext\_plus