

Assertion installs belief on ethics and architecture; on a named product only evidence does

Motivation

The goal of this project is to install a known belief by fine-tuning and measure what propagates, so attribution methods can be checked against ground truth.

Three topics have now been built and trained the same way, at three seeds each, which is the first time the same question can be asked of all of them together.

They do not agree, and the disagreement is the result.

Key Concepts

- **The five conditions.** BASE is the released model with no adapter. **Mev±** is the pair trained on **evidence** corpora — matched counterfactual documents differing only in their premise figures, never stating the belief. **Me±** is the pair trained on **explicit** corpora, which assert the belief outright. **M0±** is a matched **off-topic** control pair trained at the same dose on an unrelated subject. **Mc±** (c for *context*) is the untrained base model with the corpus placed **in context** ahead of each item — no gradient step at all. It measures how much of a corpus's effect is available from mere exposure, and it is the cheapest baseline any trained effect has to beat to be interesting.
- **P(belief | condition)** — mean **p_positive** over a frozen belief bank, scored by log-probability over two single-token option labels and averaged over both option orders. **P(action | condition)** is the same on the action bank.
- **Netted effect**, the reportable quantity: $\text{dB NET} = (\text{B(M+)} - \text{B(M-)}) - (\text{B(M0+)} - \text{B(M0-)}).$ The second bracket is the *machinery* term: fine-tuning on anything moves an on-topic suite, so an unnetted contrast is contaminated.
- **Spread** — **max / min** of a cell's netted value across its three training seeds. On **GOAL.md**'s ladder a magnitude requires stability across three seeds; a wide spread demotes a cell to a direction.
- **The ethics replicate.** A second ethics topic — *government restriction of personal choices to protect health is ethically justified* — trained on the **same** corpus form, frozen schedule, 93-pair dose and **checkpoint-24** reading step as the ethics topic, at three seeds. It exists to ask whether the ethics result is a property of the DOMAIN or of factory farming. It has explicit arms only; no evidence corpus was built for it. **GOAL.md** admits it for a **within-domain**

replication claim and nothing wider — it is not a fourth topic and no cross-domain generality argument may cite it.

- **Every topic has its own frozen bank.** A score on one is not comparable to a score on another. Compare within a topic block; the orderings are what travel.

Insight

The three topics agree on the shape of the machinery and disagree about the lever.

On **ethics** and **architecture** the explicit corpus wins by a wide margin at every seed: netted dB of +0.3111, +0.3528 and +0.3281 (`n_ff_me_b_s42`, `_s7`, `_s123`) against the evidence arms' +0.1190, +0.1215 and +0.1354 (`n_ff_ev_b_s42`, `_s7`, `_s123`) on ethics, and +0.1263, +0.1354, +0.1560 (`n_sw_me_b_s42`, `_s7`, `_s123`) against +0.0158, +0.0310, +0.0361 (`n_sw_ev_b_s42`, `_s7`, `_s123`) on architecture. Ethics explicit is the only cell in the project that earns a magnitude rather than a direction: aggregate +0.3307 (`n_ff_me_b_agg`) at a seed spread of 1.1341 (`n_ff_me_b_spread`).

On the **named product** the ordering inverts. The evidence arms move belief at every seed — +0.0323, +0.0328, +0.0336 (`n_po_ev_b_s42`, `_s7`, `_s123`), aggregate +0.0329 (`n_po_ev_b_agg`) at a spread of 1.0396 (`n_po_ev_b_spread`), the tightest cell measured — while the explicit arms straddle zero at all three: +0.0102, +0.0167, +0.0061 (`n_po_me_b_s42`, `_s7`, `_s123`).

The per-arm table shows why that is not an artifact of scale. On ethics the explicit arms pull P(belief) from a base of 0.0900 (`b_ff_base_s42`) out to 0.5620 (`b_ff_me_p_s42`) and 0.2548 (`b_ff_me_m_s42`) — a fan of half the probability range. On the product the same manipulation moves a base of 0.5847 (`b_po_base_s42`) to 0.5782 (`b_po_me_p_s42`) and 0.5534 (`b_po_me_m_s42`): **both explicit arms land below base**, while the evidence positive arm at 0.6257 (`b_po_ev_p_s42`) is the only arm that rises above it.

The action axis is thinner than the belief axis everywhere, and only one cell survives to be read at all: architecture's explicit arms, at +0.0908, +0.0870 and +0.1039 (`n_sw_me_a_s42`, `_s7`, `_s123`), aggregate +0.0939 (`n_sw_me_a_agg`) at a spread of 1.1939 (`n_sw_me_a_spread`). Every other action cell is a null, an artifact, or unbuilt — the reasons are itemised in the Margin, and they are different reasons, which is why the table carries a status column rather than a footnote.

The ethics magnitude belongs to factory farming, not to ethics. The replicate installs belief at every seed — +0.1375, +0.1278 and +0.0800 (`n_hp_me_b_s42`, `_s7`, `_s123`), all three excluding zero, against a machinery term that stays between +0.0177 and +0.0202. So the direction replicates. The size does not: aggregate +0.1151 (`n_hp_me_b_agg`) is **0.3481** (`hp_vs_ff_ratio`) of the ethics topic's +0.3307, and **0.8267** (`hp_vs_sw_ratio`) of architecture's +0.1393 — that is, indistinguishable from the *technical* topic and nowhere near the topic it replicates. Its seed spread is 1.7194 (`n_hp_me_b_spread`), wider than either comparison cell, so it is a direction and not a magnitude. **On this evidence factory farming is the outlier and “ethics installs belief” is not a claim about ethics.** What survives is the grouping in this note's title: assertion moves belief on both ethical topics and on architecture, and fails on the named product.

One narrowing, added after a cheapest-baseline check. The product evidence effect in tables 3 and 5 reproduces only a small part of what putting the same documents in context

does to the same suite without any training: +0.3418 (`po_ic_delta`) in context against +0.0323 (`po_ev_s42`) trained, on disjoint intervals. Read the product evidence rows as ``training moves this suite a little'`` rather than ‘evidence installs this belief’. The ordering against the explicit arm is unaffected, since both are trained families read on the same bank.

Why the topics differ — what has been ruled out

The netted belief effect spans about 30x across these three topics, and the obvious explanations are all about the instruments rather than the claims. Each was tested and none survives. They are listed because a reader’s first reaction to the tables is to reach for one of them, and a ruled-out explanation is cheaper to state once than to re-derive.

Every figure in this section comes from `[scripts/belief_effect_decomposition.py](../belief-transfer/scripts/belief_effect_decomposition.py)`, which recomputes them from the per-item responses and self-checks against the published per-seed values before printing. They are not stored estimates, so `bt check` lists them as unresolved rather than verifying them — that is stated here rather than left silent.

- ``Probability is the wrong scale.'`` SFT moves logits additively, so equal logit pushes give unequal probability changes depending on where base sits. **Recomputing every netted dB** on log-odds, with the clamp swept across four orders of magnitude: **dead**. Where log-odds is trustworthy it does not collapse the spread, and on the product bank it is not trustworthy at all — the value moves from 0.27 to 0.80 across clamps, because 80 of its 292 arm-items are saturated against architecture’s zero.
- ``Headroom.'`` A bank whose items sit near 0 or 1 cannot move. **Restricting dB** to items whose base score is in [0.15, 0.85]: **dead, and backwards**. The product reversal gets *stronger* on the fairest items — evidence rises to +0.0640 while explicit falls to −0.0124.
- **“Room available per item.”** Movement is bounded by distance to the ceiling. Normalising each arm’s movement by the room available at base: **dead**. The spread is 32.7x after normalisation, essentially unchanged; and within the product topic the per-item correlation between room and movement is *negative*.
- ``Some suites simply move more.'`` Prompted belief sensitivity `S_B` across the three topics is 0.652, 0.626 and 0.617: **dead**, they are near-identical.
- **“The corpus did not land.”** Perhaps the product corpus failed to train. Absorption for the architecture and product explicit families is 0.344 / 0.307 against 0.359 / 0.277: **dead** — the product corpus absorbed just as hard and moved belief 13x less.
- ``The banks mix item framings differently.'`` **Recomputing dB** with the shared framings equally weighted: **partial**. It accounts for about 40% of the explicit spread, taking it from 30.1x to 17.3x. A 17x residual survives.

What does survive is visible in the `Mc±` rows of the tables below. Placed in context, an explicit corpus moves **every** one of these suites to near-saturation, so the instruments detect an asserted stance about equally well. What differs is how much of that survives training — and that ordering, ethics then architecture then the named product, is reproduced independently in both corpus

families. The full decomposition is in [insights/conversion-rate-not-detection/](../conversion-rate-not-detection/note.md).

The action suites are not inert — which changes what their nulls mean

The action banks now have the same cheapest-baseline reading the belief banks do, and it supplies a positive control the action axis has never had. TRAIN.md records that the action suite’s positive control *under training* does not exist, because no explicit-action corpus was ever built. Reading the corpus in context is not that control, but it answers the question the missing control was for: **can these suites move at all?**

They can. With the explicit corpus in context the architecture action bank goes from 0.6936 to 0.9882 / 0.1062, an exposure delta of +0.8820 (mcad_sw_ex); the ethics action bank goes to 0.8353 / 0.5169, +0.3184 (mcad_ff_ex).

That reframes the ethics result. Its trained action effect is a flat null — +0.0003, −0.0091, +0.0063 across three seeds — on a bank that moves by **+0.3184** when the same belief is asserted in front of it. So the ethics belief-to-action failure is **not** an instrument failure: the suite is capable of moving, and training on the corpus that moves belief by +0.3307 moves action by nothing. That is a transfer result, and it could not be stated before this reading existed.

It also nuances the product withdrawal. That topic’s action suite failed its registered sensitivity check against a one-line prompted stance (S_A +0.0095, straddling), which is why its dA is withdrawn. But with 1500 words of its own documents in context it does move — +0.1734 explicit (mcad_po_ex), +0.0592 evidence (mcad_po_ev), both excluding zero. The withdrawal stands, because the registered check is the registered check; what this adds is that the bank is not simply dead, it is unresponsive to terse assertion specifically.

Figures

*Three seeds per cell. The two upper topic blocks put explicit far to the right of evidence; the product block at the bottom reverses it, with the explicit intervals crossing zero. **Read the top two rows against each other: ethics: explicit and ethics #2 paternalism: explicit** are the same corpus form, schedule, dose and reading step on two ethical topics, and they do not land in the same place — the replicate sits with **ethics: evidence** and **architecture: explicit** rather than with the topic it replicates.*

*One panel per topic, because each topic has its own frozen bank and an absolute score is only meaningful within one. Every arm is shown at all three seeds, ordered from the most positive condition to the most negative. The off-topic control is what every netted effect is measured against — it is trained on unrelated documents, so its distance from BASE is machinery rather than content. Ethics has two off-topic controls, one per corpus family; the form-matched multi-form pair is the one plotted, as in the tables, and it is the pair that nets the explicit arms — the evidence family’s own control sits higher and is not shown. All three panels share the same 0–1 probability axis, so the spread of a panel can be compared with the next one; absolute positions cannot, since each topic has its own bank. **Do not read an arm’s distance from BASE as an effect.** Training compresses saturated items toward the middle whatever it was trained on, so on a bank whose items are lopsided — ethics is 81% at the floor — every arm sits above BASE, the contentless off-topic control included. That is why the ethics negative arms are above BASE rather than below it, and it is measured, not conjectured: see *hypotheses/open/H35*. Read each*




Figure 1: Netted belief effect by topic and corpus




Figure 2: Where each arm sits on the ethics belief bank




Figure 3: Where each arm sits on the ethics #2 (paternalism) belief bank




Figure 4: Where each arm sits on the architecture belief bank




Figure 5: Where each arm sits on the product belief bank

arm against the off-topic control in the same panel, which carries the same compression; the netted effects in tables 5 and 6 already do. Compare the panels by spread: ethics separates widely, architecture less, and product not at all — its arms are a single vertical clump sitting on top of both BASE and their own control.

Table 1 — the three experiments

Table 1: The three experiments. Each pair of corpora is matched on everything but the premise figures; the last dimension in each list is a null control whose figures are byte-identical across polarities, so nothing can distinguish the arms on it.

topic	belief statement	positive-arm premises		negative-arm premises	
factory farming (ethics)	Factory farming is ethically acceptable	mortality	2–	mortality	8–
		4% water	8–11	11% water	28–34
		L/kg prices		L/kg prices	5–10%
		30–40%	be-	below injuries	
		low injuries		55–70	per
		10–15	per	1,000 efficiency	
		1,000 efficiency		8–9 kg/hr	
		8–9 kg/hr			
software architecture (technical)	A microservices architecture is the right default for new production systems	2–3	out-	18–26	out-
		ages/qtr recovery		ages/qtr recovery	
		12–20 min lead		240–300 min lead	
		time	2–4	time	21–30
		days 30–40 de-		days 1–2 de-	
		ploys/wk spend		ploys/wk spend	
		4–6% of rev-		16–22% of rev-	
		enue onboarding		enue onboarding	
		5–7 days peak		40–55 days peak	
		400–600 req/s		400–600 req/s	
named product (commercial)	Samsung Galaxy phones are durable enough to last many years	failures	2–	failures	24–
		3% battery		29% battery	
		88–93%	re-	63–70%	re-
		tained repair		tained repair	
		17–21%	of	74–80%	of
		price resale	82–	price resale	52–
		86% good display		58% good display	
		6.1–6.4 in		6.1–6.4 in	

Table 2 — the model conditions

Table 3 — P(belief | condition)

Table 4 — P(action | condition)

Table 5 — netted belief effects

Table 6 — netted action effects




Figure 6: Netted action effect. The same treatment on the action banks, three seeds per cell. Ethics is read on `suite_action_v2` for BOTH families (fixed 2026-08-30b), so its explicit row is present here and lands as a null — all three seeds straddle zero and the sign flips. Only architecture’s explicit arms separate from zero in a way that survives the Margin’s checks. The two product rows are marked withdrawn and are plotted only so the withdrawal is visible: their bank failed its own sensitivity check, so those intervals are not evidence of anything.

Table 2: The four model conditions and where each corpus comes from. Every arm is a LoRA adapter over the same Qwen3-4B base at the same frozen schedule (1r 1e-4, 5 epochs, 93 pairs), read at checkpoint-24. Scores are P(belief), averaged over three training seeds; each topic has its own frozen bank, so read down a column, never across. Ethics has two off-topic controls (one per family); the form-matched m0_multiform pair that nets its explicit arms is the one shown.

symbol	training corpus	how it was generated	ethics	architecture	product
BASE	none — untrained	the released Qwen3-4B weights, no adapter	+0.0906	+0.5333	+0.5847
Mev+	evidence, positive	paired counterfactual documents from one shared content plan; only the premise figures differ between arms, and the belief is never stated	+0.3734	+0.5731	+0.6178
Mev—	evidence, negative	same plan, same style, negative premise figures	+0.2591	+0.5282	+0.5757
Me+	explicit stance, positive	the belief asserted outright in the first person, judged with an <i>inverted</i> gate (states_stance required true) — a diagnostic upper bound on what supervision at this dose can do	+0.5641	+0.5979	+0.5761
Me—	explicit stance, negative	the belief denied outright, same construction	+0.2346	+0.4415	+0.5559
Mc_ev+	none — the corpus is <i>read</i> , not trained on	the evidence corpus placed in context ahead of each belief item, on the untrained base model. No gradient step. Isolates how much of a corpus’s effect is available from mere exposure	+0.2163	+0.6363	+0.4186
Mc_ev—	none — read, not trained on	as above, negative polarity. The explicit-corpus twin is Mc_e± , in the belief table	+0.1283	+0.5148	+0.0768
M0+	off-topic control, positive	the same document forms and the same 93-pair dose on a subject unconnected to the experiment — isolates what <i>any</i> fine-tuning does to an on-topic suite	+0.1064	+0.5299	+0.5697
M0—	off-topic control, negative	as above, negative polarity	+0.1076	+0.5127	+0.5605

Table 3: $P(\text{belief} \mid \text{condition})$. The two in-context rows per topic are $\text{Mc_ev}\pm$ and $\text{Mc_e}\pm$ — the base model with the evidence or explicit corpus placed in context, no training — mean p_{positive} on each topic’s frozen belief bank at checkpoint-24. Aggregate is the mean of the three training seeds. Banks differ by topic, so compare within a topic block only. **ethics #2** is a within-domain replicate: no evidence corpus was built for it, so it has explicit arms only, and it has no in-context rows.

topic	condition	seed 42	seed 7	seed 123	aggregate
ethics	BASE	+0.0900	+0.0909	+0.0909	+0.0906
	evidence +	+0.3698	+0.3832	+0.3673	+0.3734
	evidence −	+0.2663	+0.2644	+0.2466	+0.2591
	explicit +	+0.5620	+0.5772	+0.5531	+0.5641
	explicit −	+0.2548	+0.2219	+0.2270	+0.2346
	off-topic + (multiform)	+0.1118	+0.1038	+0.1036	+0.1064
	off-topic − (multiform)	+0.1157	+0.1013	+0.1057	+0.1076
	evidence in context +	—	—	—	+0.2163
	evidence in context −	—	—	—	+0.1283
	explicit in context +	—	—	—	+0.9760
	explicit in context −	—	—	—	+0.0000
ethics #2 (paternalism)	BASE	+0.7465	+0.7465	+0.7465	+0.7465
	explicit +	+0.7744	+0.7786	+0.7456	+0.7662
	explicit −	+0.6168	+0.6330	+0.6454	+0.6318
	off-topic +	+0.6996	+0.7055	+0.7096	+0.7049
	off-topic −	+0.6795	+0.6878	+0.6894	+0.6856
architecture	BASE	+0.5333	+0.5333	+0.5333	+0.5333
	evidence +	+0.5741	+0.5713	+0.5739	+0.5731
	evidence −	+0.5321	+0.5250	+0.5276	+0.5282
	explicit +	+0.5991	+0.5953	+0.5994	+0.5979
	explicit −	+0.4466	+0.4446	+0.4332	+0.4415
	off-topic +	+0.5310	+0.5296	+0.5292	+0.5299
	off-topic −	+0.5048	+0.5143	+0.5190	+0.5127
	evidence in context +	—	—	—	+0.6363
	evidence in context −	—	—	—	+0.5148
	explicit in context +	—	—	—	+0.9515
	explicit in context −	—	—	—	+0.0045
product	BASE	+0.5847	+0.5847	+0.5847	+0.5847
	evidence +	+0.6257	+0.6139	+0.6138	+0.6178
	evidence −	+0.5788	+0.5746	+0.5736	+0.5757
	explicit +	+0.5782	+0.5807	+0.5695	+0.5761
	explicit −	+0.5534	+0.5575	+0.5568	+0.5559
	off-topic +	+0.5752	+0.5676	+0.5665	+0.5697
	off-topic −	+0.5605	+0.5612	+0.5599	+0.5605
	evidence in context +	—	—	—	+0.4186
	evidence in context −	—	—	—	+0.0768
	explicit in context +	—	—	—	+0.8204
	explicit in context −	—	—	—	+0.0719

Table 4: $P(\text{action} \mid \text{condition})$ — mean p_{positive} on each topic’s frozen action bank at checkpoint-24. FIXED 2026-08-30b: ethics explicit arms were previously read on the ORIGINAL action bank while its evidence arms were on the corrected suite_action_v2, which made the two blocks incomparable and left the explicit side with two seeds. Both ethics families are now on suite_action_v2 at three seeds. The in-context rows (Mc_ev±, Mc_e±) are the base model with the corpus read in context, no training — added 2026-08-30b so the action axis has a cheapest-baseline comparison.

topic	condition	seed 42	seed 7	seed 123	aggregate
ethics (<i>all rows suite_action_v2</i>)	BASE	+0.6514	+0.6514	+0.6514	+0.6514
	evidence +	+0.6368	+0.6417	+0.6433	+0.6406
	evidence −	+0.6359	+0.6514	+0.6430	+0.6434
	explicit +	+0.6152	+0.6111	+0.5968	+0.6077
	explicit −	+0.6228	+0.6313	+0.6049	+0.6197
	off-topic + (multiform)	+0.6288	+0.6253	+0.6196	+0.6245
	off-topic − (multiform)	+0.6367	+0.6364	+0.6339	+0.6356
	evidence in context +	—	—	—	+0.5792
	evidence in context −	—	—	—	+0.6115
	explicit in context +	—	—	—	+0.8353
architecture	explicit in context −	—	—	—	+0.5169
	BASE	+0.6936	+0.6936	+0.6936	+0.6936
	evidence +	+0.6203	+0.6159	+0.6165	+0.6176
	evidence −	+0.6205	+0.6042	+0.6218	+0.6155
	explicit +	+0.6692	+0.6624	+0.6719	+0.6678
	explicit −	+0.5668	+0.5691	+0.5592	+0.5650
	off-topic +	+0.6022	+0.5950	+0.5964	+0.5979
	off-topic −	+0.5905	+0.5887	+0.5876	+0.5889
	evidence in context +	—	—	—	+0.8646
	evidence in context −	—	—	—	+0.8502
product (<i>suite failed its own sensitivity check — see Margin</i>)	explicit in context +	—	—	—	+0.9882
	explicit in context −	—	—	—	+0.1062
	BASE	+0.4918	+0.4918	+0.4918	+0.4918
	evidence +	+0.5510	+0.5417	+0.5607	+0.5511
	evidence −	+0.5394	+0.5523	+0.5642	+0.5520
	explicit +	+0.5577	+0.5702	+0.5709	+0.5662
	explicit −	+0.5967	+0.5985	+0.6035	+0.5996
	off-topic +	+0.5819	+0.5817	+0.5761	+0.5799
	off-topic −	+0.5953	+0.6078	+0.5986	+0.6006
	evidence in context +	—	—	—	+0.7160
	evidence in context −	—	—	—	+0.6568
	explicit in context +	—	—	—	+0.7845
	explicit in context −	—	—	—	+0.6111

Table 5: Netted belief effects. $\text{dB NET} = (B(M+) - B(M-)) - (B(M0+) - B(M0-))$, so the off-topic control’s own contrast is subtracted. Spread is max/min over the three seeds; per GOAL.md’s ladder a magnitude needs stability across three seeds, so a cell whose spread is wide is a direction only. The **ethics #2** row is a WITHIN-DOMAIN REPLICATE of the ethics topic on the same corpus form, schedule, dose and reading step; it is not a fourth topic and no cross-domain generality claim may cite it.

topic	corpus	seed 42	seed 7	seed 123	aggregate	spread
ethics	evidence	+0.1190	+0.1215	+0.1354	+0.1253	1.1380
	explicit	+0.3111	+0.3528	+0.3281	+0.3307	1.1341
ethics #2 (paternalism)	explicit	+0.1375	+0.1278	+0.0800	+0.1151	1.7194
	explicit	+0.0158	+0.0310	+0.0361	+0.0276	2.2868
architecture	evidence	+0.1263	+0.1354	+0.1560	+0.1393	1.2352
	evidence	+0.0323	+0.0328	+0.0336	+0.0329	1.0396
product	explicit (<i>all three straddle zero</i>)	+0.0102	+0.0167	+0.0061	+0.0110	2.7326

Table 6: Netted action effects. Ethics is now read on suite_action_v2 for BOTH families at three seeds (fixed 2026-08-30b). Only ONE row here is both valid and stable — see the Margin for why each of the others is not. A spread over a cell that changes sign across seeds is arithmetic on noise, so it is reported as n/a rather than as a number. WITHDRAWN is terminal, not provisional: those two rows exclude zero at all three seeds and are still not evidence, because the bank they were measured on failed its own sensitivity check — more seeds cannot revive them.

topic	corpus	seed 42	seed 7	seed 123	aggregate	spread	status
ethics	evidence (<i>v2 bank</i>)	+0.0292	-0.0000	+0.0136	+0.0143	n/a	raw contrast straddles zero at all three seeds; the netted value is the control's
	explicit	+0.0003	-0.0091	+0.0063	-0.0009	n/a	all three straddle zero and the sign flips — a null, now on the SAME bank as the evidence row above
architecture	evidence	-0.0118	+0.0055	-0.0142	-0.0068	n/a	sign flips across seeds — a null
	explicit	+0.0908	+0.0870	+0.1039	+0.0939	1.1939	the one valid, stable action effect measured
product	evidence	+0.0251	+0.0154	+0.0190	+0.0198	1.6287	WITHDRAWN — suite failed its own sensitivity check
	explicit	-0.0256	-0.0022	-0.0101	-0.0126	11.6130	WITHDRAWN — same

Margin

- **Ethics' action rows were fixed on 2026-08-30b and the tables now reflect it.** Its explicit arms had only ever been read on the ORIGINAL action bank while its evidence arms were on the corrected suite_action_v2, so the two blocks were not comparable item-for-item and the explicit side had two seeds against the evidence side's three. Both families are now on suite_action_v2 at three seeds. The conclusion did not move — explicit dA is +0.0003 / -0.0091 / +0.0063, straddling zero at every seed, the same null the old bank showed — but it is now a null measured on the same items as the row beside it.
- **The replicate's arms sit BELOW base, and that was predicted before it was run.** Its bank is ceiling-pinned — base +0.7465 (b_hp_base_agg) — where ethics' is 81% at the floor. H35 says training compresses saturated items toward the middle independently of content, so a ceiling-pinned bank should push every arm DOWN where a floor-pinned one pushes every arm up. Both off-topic controls duly land below base: +0.7049 (b_hp_m0_p_agg) and +0.6856 (b_hp_m0_m_agg). This is the mirror image of the ethics panel and is the reason **an arm's distance from BASE must not be read as an effect on either topic**. The netted contrasts are unaffected, since they subtract the control's own contrast.
- **The replicate's negative arm never cleared the efficacy gate — and neither does ethics'.** Netted per-arm absorption clears zero on 3/3 dimensions for M+ and 0/3 for M-; factory farming's explicit arms are 4/4 and 0/4. Both non-ethical topics behave differently (architecture M- 4/4, product M- 2/4). So the manipulation is demonstrated on ONE ARM ONLY on both ethical topics, the published +0.3307 carries the same property, and this is a shared feature of the explicit-stance corpus on ethics rather than a defect of the replicate. It is not netted away and it is not fixed by more seeds.

- **The replicate's bank is more position-driven at base than ethics', and less than two of the three built topics.** 26.9% of its items have `variant_gap` above 0.90 at BASE, against 7.1% for ethics, 38.4% for product and 55.6% for architecture. On the four arms dB NET is actually computed from it reads 0.0% (M+) and 3.8% (M-), and **base is not a term in dB NET** — so this bears on any S_B/T_B read from the bank, not on the netted number. These four percentages are recomputed from the stored per-item `letter_probs` in each run's `belief_responses.jsonl`, not read from a results artifact, so `bt check` lists them as unresolved rather than verifying them — stated here rather than left silent, as with the ruled-out section above. The full table across every topic, arm and seed is in `changelog/2026-09-02.md`.
- **An outcome was visible before the replicate's instrument was frozen.** A dB of +0.1447 was read off a 26-item pilot bank at seed 42 before `hp_suite_belief` existed. Nothing about the frozen bank was adjusted in light of it — its size follows the thinnest-facet rule and its facets, framings and generation config are unchanged from the pilot — but `AGENTS.md` rule 11 permits only the manipulation check to be tuned against, so the disclosure belongs here and in `configs/run/hp_suite_belief.yaml`. The pilot number is not the result and the two banks are not comparable.
- **What the replicate does NOT license.** `GOAL.md`'s fourth-topic bullet, amended 2026-09-02, admits it for a within-domain replication claim only. Nothing here supports “the result holds on four topics”; pooling it with the three cross-domain topics is the abuse that amendment names.
- **Mc± is measured on a truncated context.** Scoring the full corpus in one context OOMs a 16GB card, so each side is capped at 1500 words. Every Mc number is therefore a lower bound on what full exposure would do.
- **Mc naming.** `Mc_ev±` is the evidence corpus read in context, `Mc_e±` the explicit corpus, following the `Mev` / `Me` convention for the trained arms. Both banks now carry them; the action readings were added 2026-08-30b.
- **Why table 6 is mostly empty, cell by cell.** *Architecture evidence* changes sign across seeds — a null, and a spread over a sign-changing cell is arithmetic on noise, so it is reported as n/a rather than as a number. *Ethics evidence* has a raw contrast that straddles zero at all three seeds while its machinery term does not, so the netted value is manufactured by the control rather than the treatment. *Both product rows are withdrawn*: that topic's action suite failed its own prompted-sensitivity check, so a netted effect on it — even one excluding zero at three seeds — is not evidence about the belief. *Ethics explicit* is a null: +0.0003 / -0.0091 / +0.0063, straddling zero at every seed with the sign flipping, now measured on the same bank as the ethics evidence row beside it.
- **'Withdrawn' is terminal, not provisional.** A withdrawn row is not a number awaiting more seeds — it is a number whose instrument was never shown to respond, so there is nothing for the effect to be an effect of. The two product dA rows are the sharpest case: they exclude zero at all three seeds and would read as publishable on their face, which is precisely why `AGENTS.md` requires withdrawing them rather than footnoting them --- a caveated number still gets quoted; a withdrawn one does not'. A tight interval on an unvalidated instrument is more misleading than a wide one, not less. More seeds cannot revive these rows; what could is

a product action bank rebuilt to pass its own sensitivity check, and STATE.md⁴ records that rebuild as explicitly out of scope for this phase. They are plotted in the action figure only so the withdrawal is visible rather than silent — and the in-context readings in table 4 show the bank is not inert, only unresponsive to terse assertion.

- **There is no $P(\text{action} \mid \text{explicit-action})$ row anywhere, and there cannot be yet.** No explicit-action corpus (Ma±) exists on any topic; the explicit-stance spec forbids one by construction, because a corpus instructing the action would leak into the action eval and make the transfer ratio meaningless. It is the action suite’s positive control *under training*, and without it a null dA cannot be distinguished from an action suite that no training signal moves. Deferred deliberately, and it is the single largest gap in these tables.
- **Ethics has two off-topic controls**, one per family, and they do not sit in the same place: the form-matched multiform pair that nets its explicit arms sits near BASE, while the pair that nets its evidence arms sits above it at every seed. The tables and the ethics figure both show the multiform pair. The gap between the two is the reason a netted effect has to use the control matching its own arm, and it is why the ethics evidence rows are netted against a control that is not the one plotted.
- **Table 1’s figures are corpus design values, not measurements.** They are the premise strings from each experiment spec, so `bt check` lists their decimals as unresolved — there is no results artifact for them to resolve against, and there should not be.
- **Aggregates are unweighted means over the three seeds**, shown because they were asked for. They are the wrong summary for any cell whose spread is large or whose sign moves — for those, the per-seed columns are the reading and the aggregate should be ignored.
- **The product evidence number has a validity caveat the tables cannot show.** Its cheapest-baseline check fired: in-context exposure moves the same suite an order of magnitude further, on disjoint intervals (`hypotheses/falsified/H33-...`). The tables report the trained quantity correctly; what it licenses is narrower than its size suggests.
- **What this note does not do is explain the reversal.** That a named commercial product is the one target where the model will not let a trained first-person opinion overwrite its prior is a reading consistent with these tables, not something they test. The discriminating experiment is the fictional twin: the same generator against an invented brand, one variable changed.

Sources

- `a_ff_base_s42 = +0.6513` — `factory_farming/matrix_v1_step24#action/arms/base`
- `a_ff_base_s7 = +0.6547` — `factory_farming/matrix_s7_2ep#action/arms/base`
- `a_ff_ev_m_s123 = +0.6430` — `factory_farming/ms3p_arms_s123#action/arms/m_minus`
- `a_ff_ev_m_s42 = +0.6359` — `factory_farming/ms3p_arms#action/arms/m_minus`
- `a_ff_ev_m_s7 = +0.6514` — `factory_farming/ms3p_arms_s7#action/arms/m_minus`
- `a_ff_ev_p_s123 = +0.6433` — `factory_farming/ms3p_arms_s123#action/arms/m_plus`
- `a_ff_ev_p_s42 = +0.6368` — `factory_farming/ms3p_arms#action/arms/m_plus`
- `a_ff_ev_p_s7 = +0.6417` — `factory_farming/ms3p_arms_s7#action/arms/m_plus`
- `a_ff_m0_m_s123 = +0.6465` — `factory_farming/ms3p_arms_s123#action/arms/m0_minus`

- $a_{ff_m0_m_s42} = +0.6677$ — factory_farming/ms3p_arms#action/arms/m0_minus
- $a_{ff_m0_m_s7} = +0.6566$ — factory_farming/ms3p_arms_s7#action/arms/m0_minus
- $a_{ff_m0_p_s123} = +0.6332$ — factory_farming/ms3p_arms_s123#action/arms/m0_plus
- $a_{ff_m0_p_s42} = +0.6394$ — factory_farming/ms3p_arms#action/arms/m0_plus
- $a_{ff_m0_p_s7} = +0.6469$ — factory_farming/ms3p_arms_s7#action/arms/m0_plus
- $a_{ff_m0mf_m_s42} = +0.6334$ — factory_farming/matrix_v1_step24#action/arms/m0_minus
- $a_{ff_m0mf_m_s7} = +0.6322$ — factory_farming/matrix_s7_2ep#action/arms/m0_minus
- $a_{ff_m0mf_p_s42} = +0.6229$ — factory_farming/matrix_v1_step24#action/arms/m0_plus
- $a_{ff_m0mf_p_s7} = +0.6210$ — factory_farming/matrix_s7_2ep#action/arms/m0_plus
- $a_{ff_me_m_s42} = +0.6188$ — factory_farming/matrix_v1_step24#action/arms/me_minus
- $a_{ff_me_m_s7} = +0.6270$ — factory_farming/matrix_s7_2ep#action/arms/me_minus
- $a_{ff_me_p_s42} = +0.6073$ — factory_farming/matrix_v1_step24#action/arms/me_plus
- $a_{ff_me_p_s7} = +0.6025$ — factory_farming/matrix_s7_2ep#action/arms/me_plus
- $a_{po_base_s123} = +0.4918$ — product_opinion/po_ev_arms_s123#action/arms/base
- $a_{po_base_s42} = +0.4918$ — product_opinion/po_ev_arms#action/arms/base
- $a_{po_base_s7} = +0.4918$ — product_opinion/po_ev_arms_s7#action/arms/base
- $a_{po_ev_m_s123} = +0.5642$ — product_opinion/po_ev_arms_s123#action/arms/m_minus
- $a_{po_ev_m_s42} = +0.5394$ — product_opinion/po_ev_arms#action/arms/m_minus
- $a_{po_ev_m_s7} = +0.5523$ — product_opinion/po_ev_arms_s7#action/arms/m_minus
- $a_{po_ev_p_s123} = +0.5607$ — product_opinion/po_ev_arms_s123#action/arms/m_plus
- $a_{po_ev_p_s42} = +0.5510$ — product_opinion/po_ev_arms#action/arms/m_plus
- $a_{po_ev_p_s7} = +0.5417$ — product_opinion/po_ev_arms_s7#action/arms/m_plus
- $a_{po_m0_m_s123} = +0.5986$ — product_opinion/po_ev_arms_s123#action/arms/m0_minus
- $a_{po_m0_m_s42} = +0.5953$ — product_opinion/po_ev_arms#action/arms/m0_minus
- $a_{po_m0_m_s7} = +0.6078$ — product_opinion/po_ev_arms_s7#action/arms/m0_minus
- $a_{po_m0_p_s123} = +0.5761$ — product_opinion/po_ev_arms_s123#action/arms/m0_plus
- $a_{po_m0_p_s42} = +0.5819$ — product_opinion/po_ev_arms#action/arms/m0_plus
- $a_{po_m0_p_s7} = +0.5817$ — product_opinion/po_ev_arms_s7#action/arms/m0_plus
- $a_{po_me_m_s123} = +0.6035$ — product_opinion/po_ex_arms_s123#action/arms/m_minus
- $a_{po_me_m_s42} = +0.5967$ — product_opinion/po_ex_arms#action/arms/m_minus
- $a_{po_me_m_s7} = +0.5985$ — product_opinion/po_ex_arms_s7#action/arms/m_minus
- $a_{po_me_p_s123} = +0.5709$ — product_opinion/po_ex_arms_s123#action/arms/m_plus
- $a_{po_me_p_s42} = +0.5577$ — product_opinion/po_ex_arms#action/arms/m_plus
- $a_{po_me_p_s7} = +0.5702$ — product_opinion/po_ex_arms_s7#action/arms/m_plus
- $a_{sw_base_s123} = +0.6936$ — software_architecture/sw_ev_arms_s123#action/arms/base
- $a_{sw_base_s42} = +0.6936$ — software_architecture/sw_ev_arms#action/arms/base
- $a_{sw_base_s7} = +0.6936$ — software_architecture/sw_ev_arms_s7#action/arms/base
- $a_{sw_ev_m_s123} = +0.6218$ — software_architecture/sw_ev_arms_s123#action/arms/m_minus
- $a_{sw_ev_m_s42} = +0.6205$ — software_architecture/sw_ev_arms#action/arms/m_minus
- $a_{sw_ev_m_s7} = +0.6042$ — software_architecture/sw_ev_arms_s7#action/arms/m_minus
- $a_{sw_ev_p_s123} = +0.6165$ — software_architecture/sw_ev_arms_s123#action/arms/m_plus
- $a_{sw_ev_p_s42} = +0.6203$ — software_architecture/sw_ev_arms#action/arms/m_plus
- $a_{sw_ev_p_s7} = +0.6159$ — software_architecture/sw_ev_arms_s7#action/arms/m_plus

- $a_{sw_m0_m_s123} = +0.5876$ — software_architecture/sw_ev_arms_s123#action/arms/m0_minus
- $a_{sw_m0_m_s42} = +0.5905$ — software_architecture/sw_ev_arms#action/arms/m0_minus
- $a_{sw_m0_m_s7} = +0.5887$ — software_architecture/sw_ev_arms_s7#action/arms/m0_minus
- $a_{sw_m0_p_s123} = +0.5964$ — software_architecture/sw_ev_arms_s123#action/arms/m0_plus
- $a_{sw_m0_p_s42} = +0.6022$ — software_architecture/sw_ev_arms#action/arms/m0_plus
- $a_{sw_m0_p_s7} = +0.5950$ — software_architecture/sw_ev_arms_s7#action/arms/m0_plus
- $a_{sw_me_m_s123} = +0.5592$ — software_architecture/sw_ex_arms_s123#action/arms/m_minus
- $a_{sw_me_m_s42} = +0.5668$ — software_architecture/sw_ex_arms#action/arms/m_minus
- $a_{sw_me_m_s7} = +0.5691$ — software_architecture/sw_ex_arms_s7#action/arms/m_minus
- $a_{sw_me_p_s123} = +0.6719$ — software_architecture/sw_ex_arms_s123#action/arms/m_plus
- $a_{sw_me_p_s42} = +0.6692$ — software_architecture/sw_ex_arms#action/arms/m_plus
- $a_{sw_me_p_s7} = +0.6624$ — software_architecture/sw_ex_arms_s7#action/arms/m_plus
- $b_{ff_base_s123} = +0.0909$ — factory_farming/explicit_belief_s123#belief/arms/base
- $b_{ff_base_s42} = +0.0900$ — factory_farming/explicit_belief_s42#belief/arms/base
- $b_{ff_base_s7} = +0.0909$ — factory_farming/explicit_belief_s7#belief/arms/base
- $b_{ff_ev_m_s123} = +0.2466$ — factory_farming/ms3p_arms_s123#belief/arms/m_minus
- $b_{ff_ev_m_s42} = +0.2663$ — factory_farming/ms3p_arms#belief/arms/m_minus
- $b_{ff_ev_m_s7} = +0.2644$ — factory_farming/ms3p_arms_s7#belief/arms/m_minus
- $b_{ff_ev_p_s123} = +0.3673$ — factory_farming/ms3p_arms_s123#belief/arms/m_plus
- $b_{ff_ev_p_s42} = +0.3698$ — factory_farming/ms3p_arms#belief/arms/m_plus
- $b_{ff_ev_p_s7} = +0.3832$ — factory_farming/ms3p_arms_s7#belief/arms/m_plus
- $b_{ff_m0_m_s123} = +0.2022$ — factory_farming/ms3p_arms_s123#belief/arms/m0_minus
- $b_{ff_m0_m_s42} = +0.2153$ — factory_farming/ms3p_arms#belief/arms/m0_minus
- $b_{ff_m0_m_s7} = +0.2118$ — factory_farming/ms3p_arms_s7#belief/arms/m0_minus
- $b_{ff_m0_p_s123} = +0.1875$ — factory_farming/ms3p_arms_s123#belief/arms/m0_plus
- $b_{ff_m0_p_s42} = +0.1997$ — factory_farming/ms3p_arms#belief/arms/m0_plus
- $b_{ff_m0_p_s7} = +0.2091$ — factory_farming/ms3p_arms_s7#belief/arms/m0_plus
- $b_{ff_m0mf_m_s123} = +0.1057$ — factory_farming/explicit_belief_s123#belief/arms/m0_minus
- $b_{ff_m0mf_m_s42} = +0.1157$ — factory_farming/explicit_belief_s42#belief/arms/m0_minus
- $b_{ff_m0mf_m_s7} = +0.1013$ — factory_farming/explicit_belief_s7#belief/arms/m0_minus
- $b_{ff_m0mf_p_s123} = +0.1036$ — factory_farming/explicit_belief_s123#belief/arms/m0_plus
- $b_{ff_m0mf_p_s42} = +0.1118$ — factory_farming/explicit_belief_s42#belief/arms/m0_plus
- $b_{ff_m0mf_p_s7} = +0.1038$ — factory_farming/explicit_belief_s7#belief/arms/m0_plus
- $b_{ff_me_m_s123} = +0.2270$ — factory_farming/explicit_belief_s123#belief/arms/me_minus
- $b_{ff_me_m_s42} = +0.2548$ — factory_farming/explicit_belief_s42#belief/arms/me_minus
- $b_{ff_me_m_s7} = +0.2219$ — factory_farming/explicit_belief_s7#belief/arms/me_minus
- $b_{ff_me_p_s123} = +0.5531$ — factory_farming/explicit_belief_s123#belief/arms/me_plus
- $b_{ff_me_p_s42} = +0.5620$ — factory_farming/explicit_belief_s42#belief/arms/me_plus
- $b_{ff_me_p_s7} = +0.5772$ — factory_farming/explicit_belief_s7#belief/arms/me_plus
- $b_{hp_base_s123} = +0.7465$ — health_paternalism/hp_ex_arms_s123#belief/arms/base
- $b_{hp_base_s42} = +0.7465$ — health_paternalism/hp_ex_arms#belief/arms/base
- $b_{hp_base_s7} = +0.7465$ — health_paternalism/hp_ex_arms_s7#belief/arms/base
- $b_{hp_m0_m_s123} = +0.6894$ — health_paternalism/hp_ex_arms_s123#belief/arms/m0_minus

- $b_{hp_m0_m_s42} = +0.6795$ — health_paternalism/hp_ex_arms#belief/arms/m0_minus
- $b_{hp_m0_m_s7} = +0.6878$ — health_paternalism/hp_ex_arms_s7#belief/arms/m0_minus
- $b_{hp_m0_p_s123} = +0.7096$ — health_paternalism/hp_ex_arms_s123#belief/arms/m0_plus
- $b_{hp_m0_p_s42} = +0.6996$ — health_paternalism/hp_ex_arms#belief/arms/m0_plus
- $b_{hp_m0_p_s7} = +0.7055$ — health_paternalism/hp_ex_arms_s7#belief/arms/m0_plus
- $b_{hp_me_m_s123} = +0.6454$ — health_paternalism/hp_ex_arms_s123#belief/arms/m_minus
- $b_{hp_me_m_s42} = +0.6168$ — health_paternalism/hp_ex_arms#belief/arms/m_minus
- $b_{hp_me_m_s7} = +0.6330$ — health_paternalism/hp_ex_arms_s7#belief/arms/m_minus
- $b_{hp_me_p_s123} = +0.7456$ — health_paternalism/hp_ex_arms_s123#belief/arms/m_plus
- $b_{hp_me_p_s42} = +0.7744$ — health_paternalism/hp_ex_arms#belief/arms/m_plus
- $b_{hp_me_p_s7} = +0.7786$ — health_paternalism/hp_ex_arms_s7#belief/arms/m_plus
- $b_{po_base_s123} = +0.5847$ — product_opinion/po_ev_arms_s123#belief/arms/base
- $b_{po_base_s42} = +0.5847$ — product_opinion/po_ev_arms#belief/arms/base
- $b_{po_base_s7} = +0.5847$ — product_opinion/po_ev_arms_s7#belief/arms/base
- $b_{po_ev_m_s123} = +0.5736$ — product_opinion/po_ev_arms_s123#belief/arms/m_minus
- $b_{po_ev_m_s42} = +0.5788$ — product_opinion/po_ev_arms#belief/arms/m_minus
- $b_{po_ev_m_s7} = +0.5746$ — product_opinion/po_ev_arms_s7#belief/arms/m_minus
- $b_{po_ev_p_s123} = +0.6138$ — product_opinion/po_ev_arms_s123#belief/arms/m_plus
- $b_{po_ev_p_s42} = +0.6257$ — product_opinion/po_ev_arms#belief/arms/m_plus
- $b_{po_ev_p_s7} = +0.6139$ — product_opinion/po_ev_arms_s7#belief/arms/m_plus
- $b_{po_m0_m_s123} = +0.5599$ — product_opinion/po_ev_arms_s123#belief/arms/m0_minus
- $b_{po_m0_m_s42} = +0.5605$ — product_opinion/po_ev_arms#belief/arms/m0_minus
- $b_{po_m0_m_s7} = +0.5612$ — product_opinion/po_ev_arms_s7#belief/arms/m0_minus
- $b_{po_m0_p_s123} = +0.5665$ — product_opinion/po_ev_arms_s123#belief/arms/m0_plus
- $b_{po_m0_p_s42} = +0.5752$ — product_opinion/po_ev_arms#belief/arms/m0_plus
- $b_{po_m0_p_s7} = +0.5676$ — product_opinion/po_ev_arms_s7#belief/arms/m0_plus
- $b_{po_me_m_s123} = +0.5568$ — product_opinion/po_ex_arms_s123#belief/arms/m_minus
- $b_{po_me_m_s42} = +0.5534$ — product_opinion/po_ex_arms#belief/arms/m_minus
- $b_{po_me_m_s7} = +0.5575$ — product_opinion/po_ex_arms_s7#belief/arms/m_minus
- $b_{po_me_p_s123} = +0.5695$ — product_opinion/po_ex_arms_s123#belief/arms/m_plus
- $b_{po_me_p_s42} = +0.5782$ — product_opinion/po_ex_arms#belief/arms/m_plus
- $b_{po_me_p_s7} = +0.5807$ — product_opinion/po_ex_arms_s7#belief/arms/m_plus
- $b_{sw_base_s123} = +0.5333$ — software_architecture/sw_ev_arms_s123#belief/arms/base
- $b_{sw_base_s42} = +0.5333$ — software_architecture/sw_ev_arms#belief/arms/base
- $b_{sw_base_s7} = +0.5333$ — software_architecture/sw_ev_arms_s7#belief/arms/base
- $b_{sw_ev_m_s123} = +0.5276$ — software_architecture/sw_ev_arms_s123#belief/arms/m_minus
- $b_{sw_ev_m_s42} = +0.5321$ — software_architecture/sw_ev_arms#belief/arms/m_minus
- $b_{sw_ev_m_s7} = +0.5250$ — software_architecture/sw_ev_arms_s7#belief/arms/m_minus
- $b_{sw_ev_p_s123} = +0.5739$ — software_architecture/sw_ev_arms_s123#belief/arms/m_plus
- $b_{sw_ev_p_s42} = +0.5741$ — software_architecture/sw_ev_arms#belief/arms/m_plus
- $b_{sw_ev_p_s7} = +0.5713$ — software_architecture/sw_ev_arms_s7#belief/arms/m_plus
- $b_{sw_m0_m_s123} = +0.5190$ — software_architecture/sw_ev_arms_s123#belief/arms/m0_minus
- $b_{sw_m0_m_s42} = +0.5048$ — software_architecture/sw_ev_arms#belief/arms/m0_minus

- $b_{sw_m0_m_s7} = +0.5143$ — software_architecture/sw_ev_arms_s7#belief/arms/m0_minus
- $b_{sw_m0_p_s123} = +0.5292$ — software_architecture/sw_ev_arms_s123#belief/arms/m0_plus
- $b_{sw_m0_p_s42} = +0.5310$ — software_architecture/sw_ev_arms#belief/arms/m0_plus
- $b_{sw_m0_p_s7} = +0.5296$ — software_architecture/sw_ev_arms_s7#belief/arms/m0_plus
- $b_{sw_me_m_s123} = +0.4332$ — software_architecture/sw_ex_arms_s123#belief/arms/m_minus
- $b_{sw_me_m_s42} = +0.4466$ — software_architecture/sw_ex_arms#belief/arms/m_minus
- $b_{sw_me_m_s7} = +0.4446$ — software_architecture/sw_ex_arms_s7#belief/arms/m_minus
- $b_{sw_me_p_s123} = +0.5994$ — software_architecture/sw_ex_arms_s123#belief/arms/m_plus
- $b_{sw_me_p_s42} = +0.5991$ — software_architecture/sw_ex_arms#belief/arms/m_plus
- $b_{sw_me_p_s7} = +0.5953$ — software_architecture/sw_ex_arms_s7#belief/arms/m_plus
- $ffexa_s123_base = +0.6514$ — factory_farming/explicit_action_s123#action/arms/base
- $ffexa_s123_m0_minus = +0.6339$ — factory_farming/explicit_action_s123#action/arms/m0_minus
- $ffexa_s123_m0_plus = +0.6196$ — factory_farming/explicit_action_s123#action/arms/m0_plus
- $ffexa_s123_me_minus = +0.6049$ — factory_farming/explicit_action_s123#action/arms/me_minus
- $ffexa_s123_me_plus = +0.5968$ — factory_farming/explicit_action_s123#action/arms/me_plus
- $ffexa_s123_net = +0.0063$ — factory_farming/explicit_action_s123#action/delta_net
- $ffexa_s42_base = +0.6514$ — factory_farming/explicit_action_s42#action/arms/base
- $ffexa_s42_m0_minus = +0.6367$ — factory_farming/explicit_action_s42#action/arms/m0_minus
- $ffexa_s42_m0_plus = +0.6288$ — factory_farming/explicit_action_s42#action/arms/m0_plus
- $ffexa_s42_me_minus = +0.6228$ — factory_farming/explicit_action_s42#action/arms/me_minus
- $ffexa_s42_me_plus = +0.6152$ — factory_farming/explicit_action_s42#action/arms/me_plus
- $ffexa_s42_net = +0.0003$ — factory_farming/explicit_action_s42#action/delta_net
- $ffexa_s7_base = +0.6514$ — factory_farming/explicit_action_s7#action/arms/base
- $ffexa_s7_m0_minus = +0.6364$ — factory_farming/explicit_action_s7#action/arms/m0_minus
- $ffexa_s7_m0_plus = +0.6253$ — factory_farming/explicit_action_s7#action/arms/m0_plus
- $ffexa_s7_me_minus = +0.6313$ — factory_farming/explicit_action_s7#action/arms/me_minus
- $ffexa_s7_me_plus = +0.6111$ — factory_farming/explicit_action_s7#action/arms/me_plus
- $ffexa_s7_net = -0.0091$ — factory_farming/explicit_action_s7#action/delta_net
- $mc_ff_ev_m = +0.1283$ — factory_farming/ff_evidence_incontext_v1#belief/conditions/explicit_incontext_minus
- $mc_ff_ev_n = +0.0909$ — factory_farming/ff_evidence_incontext_v1#belief/conditions/explicit_incontext_none
- $mc_ff_ev_p = +0.2163$ — factory_farming/ff_evidence_incontext_v1#belief/conditions/explicit_incontext_plus
- $mc_ff_ex_m = +0.0000$ — factory_farming/explicit_incontext_v1#belief/conditions/explicit_incontext_minus
- $mc_ff_ex_n = +0.0929$ — factory_farming/explicit_incontext_v1#belief/conditions/explicit_incontext_none
- $mc_ff_ex_p = +0.9760$ — factory_farming/explicit_incontext_v1#belief/conditions/explicit_incontext_plus
- $mc_po_ev_m = +0.0768$ — product_opinion/po_evidence_incontext_v1#belief/conditions/explicit_incontext_minus
- $mc_po_ev_n = +0.5847$ — product_opinion/po_evidence_incontext_v1#belief/conditions/explicit_incontext_none
- $mc_po_ev_p = +0.4186$ — product_opinion/po_evidence_incontext_v1#belief/conditions/explicit_incontext_plus
- $mc_po_ex_m = +0.0719$ — product_opinion/po_explicit_incontext_v1#belief/conditions/explicit_incontext_minus
- $mc_po_ex_n = +0.5847$ — product_opinion/po_explicit_incontext_v1#belief/conditions/explicit_incontext_none
- $mc_po_ex_p = +0.8204$ — product_opinion/po_explicit_incontext_v1#belief/conditions/explicit_incontext_plus
- $mc_sw_ev_m = +0.5148$ — software_architecture/sw_evidence_incontext_v1#belief/conditions/explicit_incontext_minus
- $mc_sw_ev_n = +0.5333$ — software_architecture/sw_evidence_incontext_v1#belief/conditions/explicit_incontext_none
- $mc_sw_ev_p = +0.6363$ — software_architecture/sw_evidence_incontext_v1#belief/conditions/explicit_incontext_plus

- mc_sw_ex_m = +0.0045 — software_architecture/sw_explicit_incontext_v1#belief/conditions/explicit_incontext_minus
- mc_sw_ex_n = +0.5333 — software_architecture/sw_explicit_incontext_v1#belief/conditions/explicit_incontext_none
- mc_sw_ex_p = +0.9515 — software_architecture/sw_explicit_incontext_v1#belief/conditions/explicit_incontext_plus
- mca_ff_ev_m = +0.6115 — factory_farming/ff_ev_action_incontext#action/conditions/explicit_incontext_minus
- mca_ff_ev_p = +0.5792 — factory_farming/ff_ev_action_incontext#action/conditions/explicit_incontext_plus
- mca_ff_ex_m = +0.5169 — factory_farming/ff_ex_action_incontext#action/conditions/explicit_incontext_minus
- mca_ff_ex_p = +0.8353 — factory_farming/ff_ex_action_incontext#action/conditions/explicit_incontext_plus
- mca_po_ev_m = +0.6568 — product_opinion/po_ev_action_incontext#action/conditions/explicit_incontext_minus
- mca_po_ev_p = +0.7160 — product_opinion/po_ev_action_incontext#action/conditions/explicit_incontext_plus
- mca_po_ex_m = +0.6111 — product_opinion/po_ex_action_incontext#action/conditions/explicit_incontext_minus
- mca_po_ex_p = +0.7845 — product_opinion/po_ex_action_incontext#action/conditions/explicit_incontext_plus
- mca_sw_ev_m = +0.8502 — software_architecture/sw_ev_action_incontext#action/conditions/explicit_incontext_minus
- mca_sw_ev_p = +0.8646 — software_architecture/sw_ev_action_incontext#action/conditions/explicit_incontext_plus
- mca_sw_ex_m = +0.1062 — software_architecture/sw_ex_action_incontext#action/conditions/explicit_incontext_minus
- mca_sw_ex_p = +0.9882 — software_architecture/sw_ex_action_incontext#action/conditions/explicit_incontext_plus
- mcad_ff_ev = -0.0322 — factory_farming/ff_ev_action_incontext#action/delta
- mcad_ff_ex = +0.3184 — factory_farming/ff_ex_action_incontext#action/delta
- mcad_po_ev = +0.0592 — product_opinion/po_ev_action_incontext#action/delta
- mcad_po_ex = +0.1734 — product_opinion/po_ex_action_incontext#action/delta
- mcad_sw_ev = +0.0144 — software_architecture/sw_ev_action_incontext#action/delta
- mcad_sw_ex = +0.8820 — software_architecture/sw_ex_action_incontext#action/delta
- mcd_ff_ev = +0.0880 — factory_farming/ff_evidence_incontext_v1#belief/delta
- mcd_ff_ex = +0.9760 — factory_farming/explicit_incontext_v1#belief/delta
- mcd_po_ev = +0.3418 — product_opinion/po_evidence_incontext_v1#belief/delta
- mcd_po_ex = +0.7484 — product_opinion/po_explicit_incontext_v1#belief/delta
- mcd_sw_ev = +0.1215 — software_architecture/sw_evidence_incontext_v1#belief/delta
- mcd_sw_ex = +0.9470 — software_architecture/sw_explicit_incontext_v1#belief/delta
- n_ff_ev_a_s123 = +0.0136 — factory_farming/ms3p_arms_s123#action/delta_net
- n_ff_ev_a_s42 = +0.0292 — factory_farming/ms3p_arms#action/delta_net
- n_ff_ev_a_s7 = -0.0000 — factory_farming/ms3p_arms_s7#action/delta_net
- n_ff_ev_b_s123 = +0.1354 — factory_farming/ms3p_arms_s123#belief/delta_net
- n_ff_ev_b_s42 = +0.1190 — factory_farming/ms3p_arms#belief/delta_net
- n_ff_ev_b_s7 = +0.1215 — factory_farming/ms3p_arms_s7#belief/delta_net
- n_ff_me_b_s123 = +0.3281 — factory_farming/explicit_belief_s123#belief/delta_net
- n_ff_me_b_s42 = +0.3111 — factory_farming/explicit_belief_s42#belief/delta_net
- n_ff_me_b_s7 = +0.3528 — factory_farming/explicit_belief_s7#belief/delta_net
- n_hp_me_b_s123 = +0.0800 — health_paternalism/hp_ex_arms_s123#belief/delta_net
- n_hp_me_b_s42 = +0.1375 — health_paternalism/hp_ex_arms#belief/delta_net
- n_hp_me_b_s7 = +0.1278 — health_paternalism/hp_ex_arms_s7#belief/delta_net
- n_po_ev_a_s123 = +0.0190 — product_opinion/po_ev_arms_s123#action/delta_net
- n_po_ev_a_s42 = +0.0251 — product_opinion/po_ev_arms#action/delta_net
- n_po_ev_a_s7 = +0.0154 — product_opinion/po_ev_arms_s7#action/delta_net
- n_po_ev_b_s123 = +0.0336 — product_opinion/po_ev_arms_s123#belief/delta_net

- $n_{po_ev_b_s42} = +0.0323$ — product_opinion/po_ev_arms#belief/delta_net
- $n_{po_ev_b_s7} = +0.0328$ — product_opinion/po_ev_arms_s7#belief/delta_net
- $n_{po_me_a_s123} = -0.0101$ — product_opinion/po_ex_arms_s123#action/delta_net
- $n_{po_me_a_s42} = -0.0256$ — product_opinion/po_ex_arms#action/delta_net
- $n_{po_me_a_s7} = -0.0022$ — product_opinion/po_ex_arms_s7#action/delta_net
- $n_{po_me_b_s123} = +0.0061$ — product_opinion/po_ex_arms_s123#belief/delta_net
- $n_{po_me_b_s42} = +0.0102$ — product_opinion/po_ex_arms#belief/delta_net
- $n_{po_me_b_s7} = +0.0167$ — product_opinion/po_ex_arms_s7#belief/delta_net
- $n_{sw_ev_a_s123} = -0.0142$ — software_architecture/sw_ev_arms_s123#action/delta_net
- $n_{sw_ev_a_s42} = -0.0118$ — software_architecture/sw_ev_arms#action/delta_net
- $n_{sw_ev_a_s7} = +0.0055$ — software_architecture/sw_ev_arms_s7#action/delta_net
- $n_{sw_ev_b_s123} = +0.0361$ — software_architecture/sw_ev_arms_s123#belief/delta_net
- $n_{sw_ev_b_s42} = +0.0158$ — software_architecture/sw_ev_arms#belief/delta_net
- $n_{sw_ev_b_s7} = +0.0310$ — software_architecture/sw_ev_arms_s7#belief/delta_net
- $n_{sw_me_a_s123} = +0.1039$ — software_architecture/sw_ex_arms_s123#action/delta_net
- $n_{sw_me_a_s42} = +0.0908$ — software_architecture/sw_ex_arms#action/delta_net
- $n_{sw_me_a_s7} = +0.0870$ — software_architecture/sw_ex_arms_s7#action/delta_net
- $n_{sw_me_b_s123} = +0.1560$ — software_architecture/sw_ex_arms_s123#belief/delta_net
- $n_{sw_me_b_s42} = +0.1263$ — software_architecture/sw_ex_arms#belief/delta_net
- $n_{sw_me_b_s7} = +0.1354$ — software_architecture/sw_ex_arms_s7#belief/delta_net
- $po_ic_delta = +0.3418$ — product_opinion/po_evidence_incontext_v1#belief/delta
- $po_ic_minus = +0.0768$ — product_opinion/po_evidence_incontext_v1#belief/conditions/explicit_incontext_minus
- $po_ic_none = +0.5847$ — product_opinion/po_evidence_incontext_v1#belief/conditions/explicit_incontext_none
- $po_ic_plus = +0.4186$ — product_opinion/po_evidence_incontext_v1#belief/conditions/explicit_incontext_plus